

云原生WASM

刘振伟



云原生WASM

In Scope: 云原生对WASM的支持

Out Of Scope: WASM自身的相关概念和原理等

WebAssembly(WASM)



Solomon Hykes @solomonstre · Mar 28, 2019

If WASM+WASI existed in 2008, we wouldn't have needed to created Docker. That's how important it is. Webassembly on the server is the future of computing. A standardized system interface was the missing link. Let's hope WASI is up to the task!



Solomon Hykes @solomonstre · Mar 28, 2019

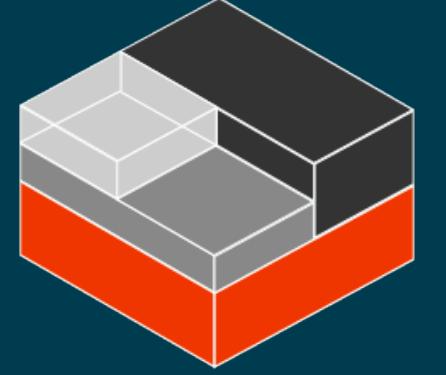
"So will wasm replace Docker?" No, but imagine a future where Docker runs linux containers, windows containers and wasm containers side by side. Over time wasm might become the most popular container type. Docker will love them all equally, and run it all :)

WASM 正在被云原生接纳

The screenshot shows the official Docker website. At the top, there's a purple header bar with the text "Wasm is a fast, light alternative to Linux containers – try it out today with the Docker+Wasm Beta." Below the header, the Docker+Wasm logo is highlighted with a red border. The main heading "Develop faster. Run anywhere." is prominently displayed in large white text. Below it, a subtext states "The most-loved Tool in Stack Overflow's 2022 Developer Survey." A blue button labeled "Download Docker Desktop" with an "Intel Chip" badge is visible. The overall background is dark blue.

Docker

The screenshot shows a GitHub commit history for the "Crun" project. The first commit is by "giuseppe" on "Dec 22, 2021". It has two commits: "1.4" and "3daded0". The commit "1.4" is highlighted with a red border. The commit message for "1.4" is "wasm: support for running on kubernetes with containerd.". The commit message for "3daded0" is "linux: add support for recursive mount options. e.g. it is possible to specify "rro" to make the mount read-only recursively.". The commit "1.4" also includes other items like "add support for idmapped mounts through a new mount option "idmap"" and "improve detection of /dev target". The commit "3daded0" includes items like "now crun exec uses CLONE_INTO_CGROUP on supported kernels when using cgroup v2" and "attempt the chdir again with the specified user if it failed before changing credentials".

**WASM****VS****Linux Container**

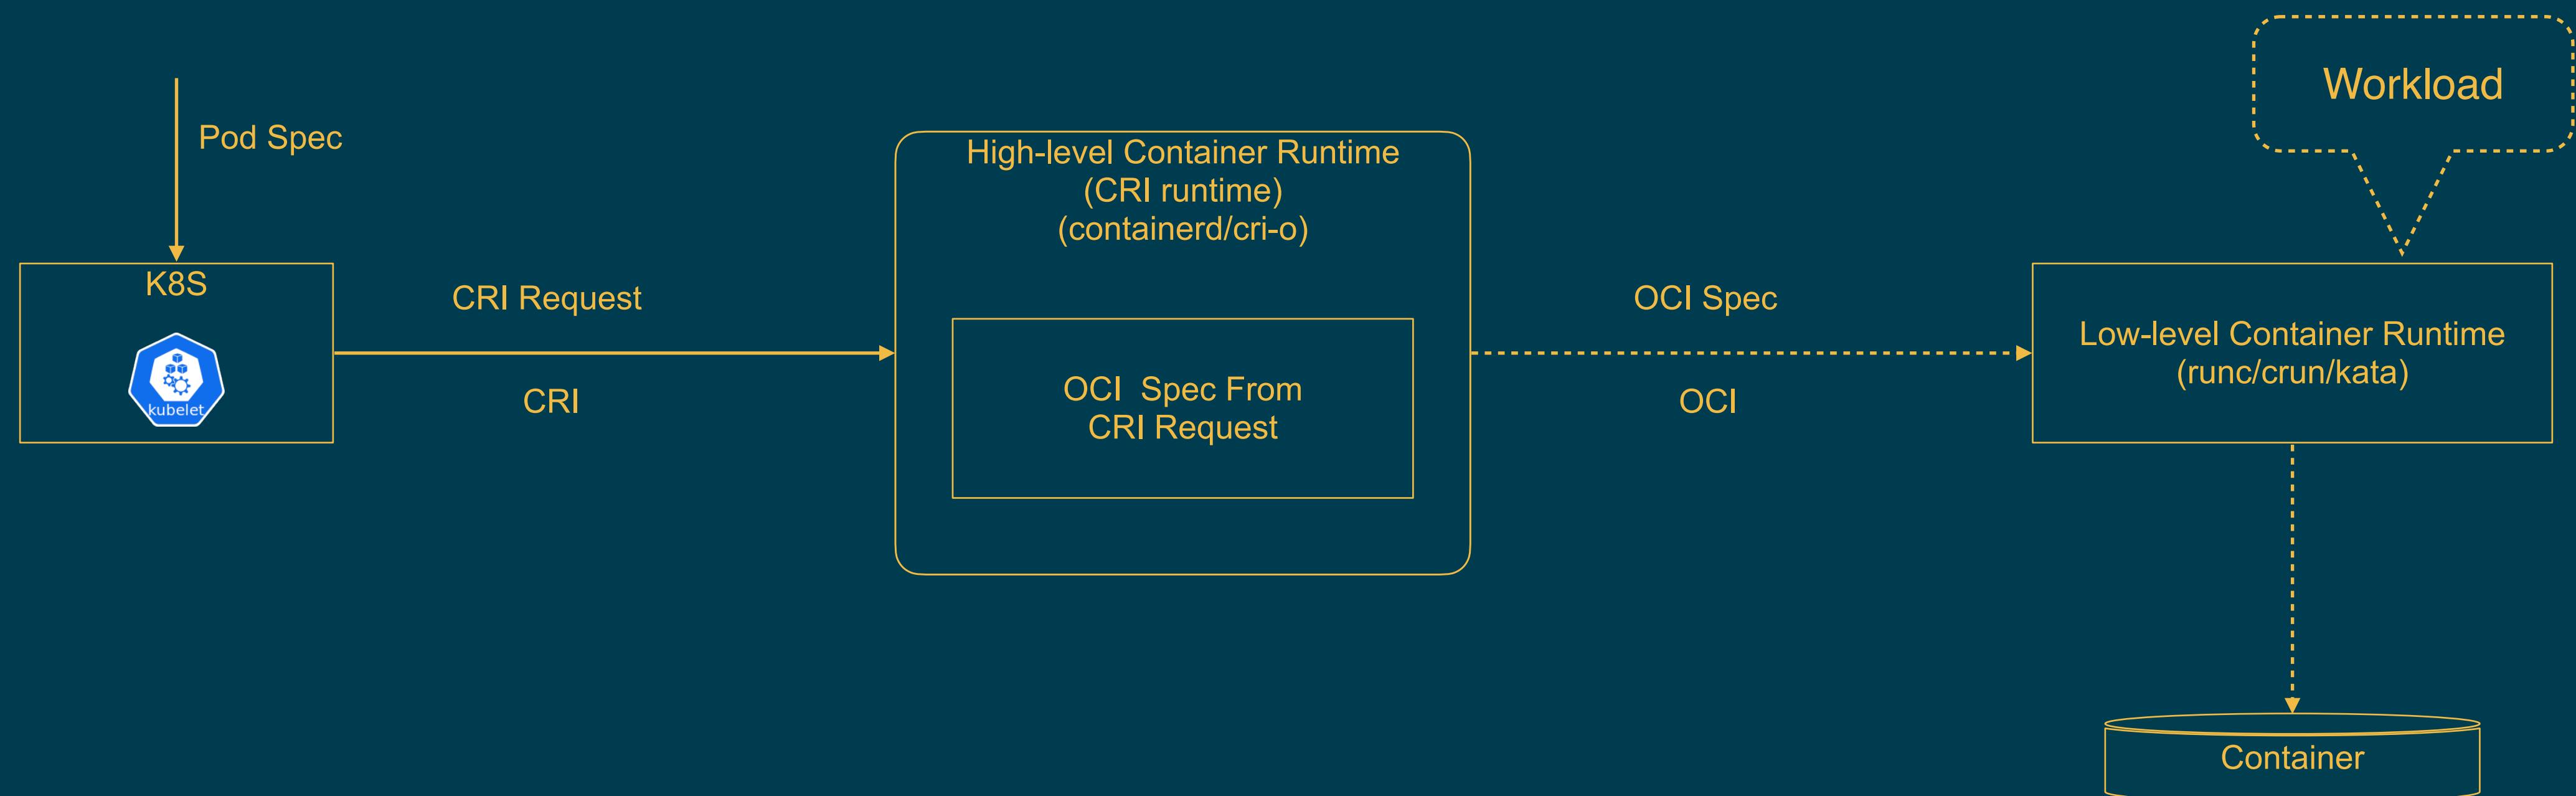
A container emulates a private operating system.

A WebAssembly instance emulates a process.

— Taylor McMullen, CTO, Fastly

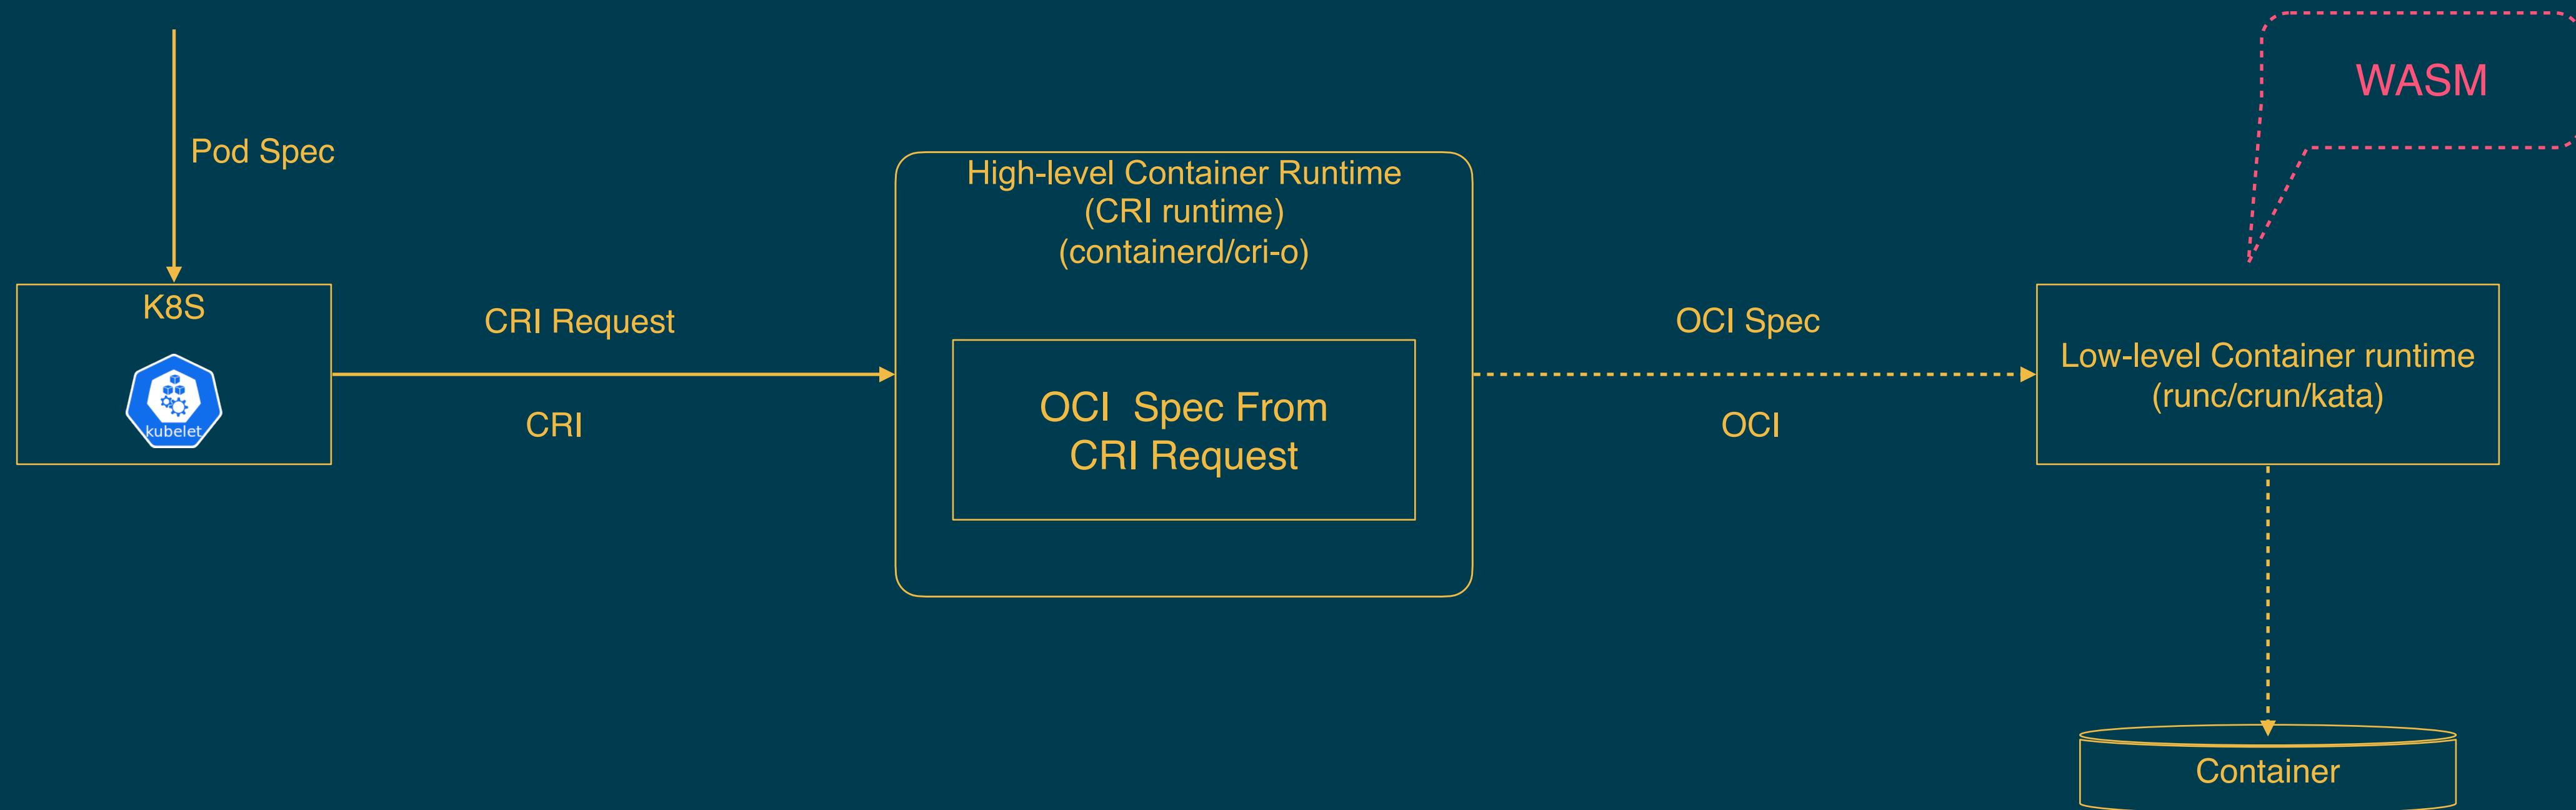
Wasm**Performance****Container****Several MBs****DISK FOOTPRINT****Several GBs****milliseconds****STARTUP PERFORMANCE****seconds****AOT is within 10% of native client****RUNTIME PERFORMANCE****10% to 20% loss to native client**

Workload 创建流程



关于docker的接入： kubelet—> cri-dockerd—>dockerd—>containerd—>.....

在Low-level Container Runtime中支持WASM Workload



通过扩展Low-level Container的能力来支持WASM Workload

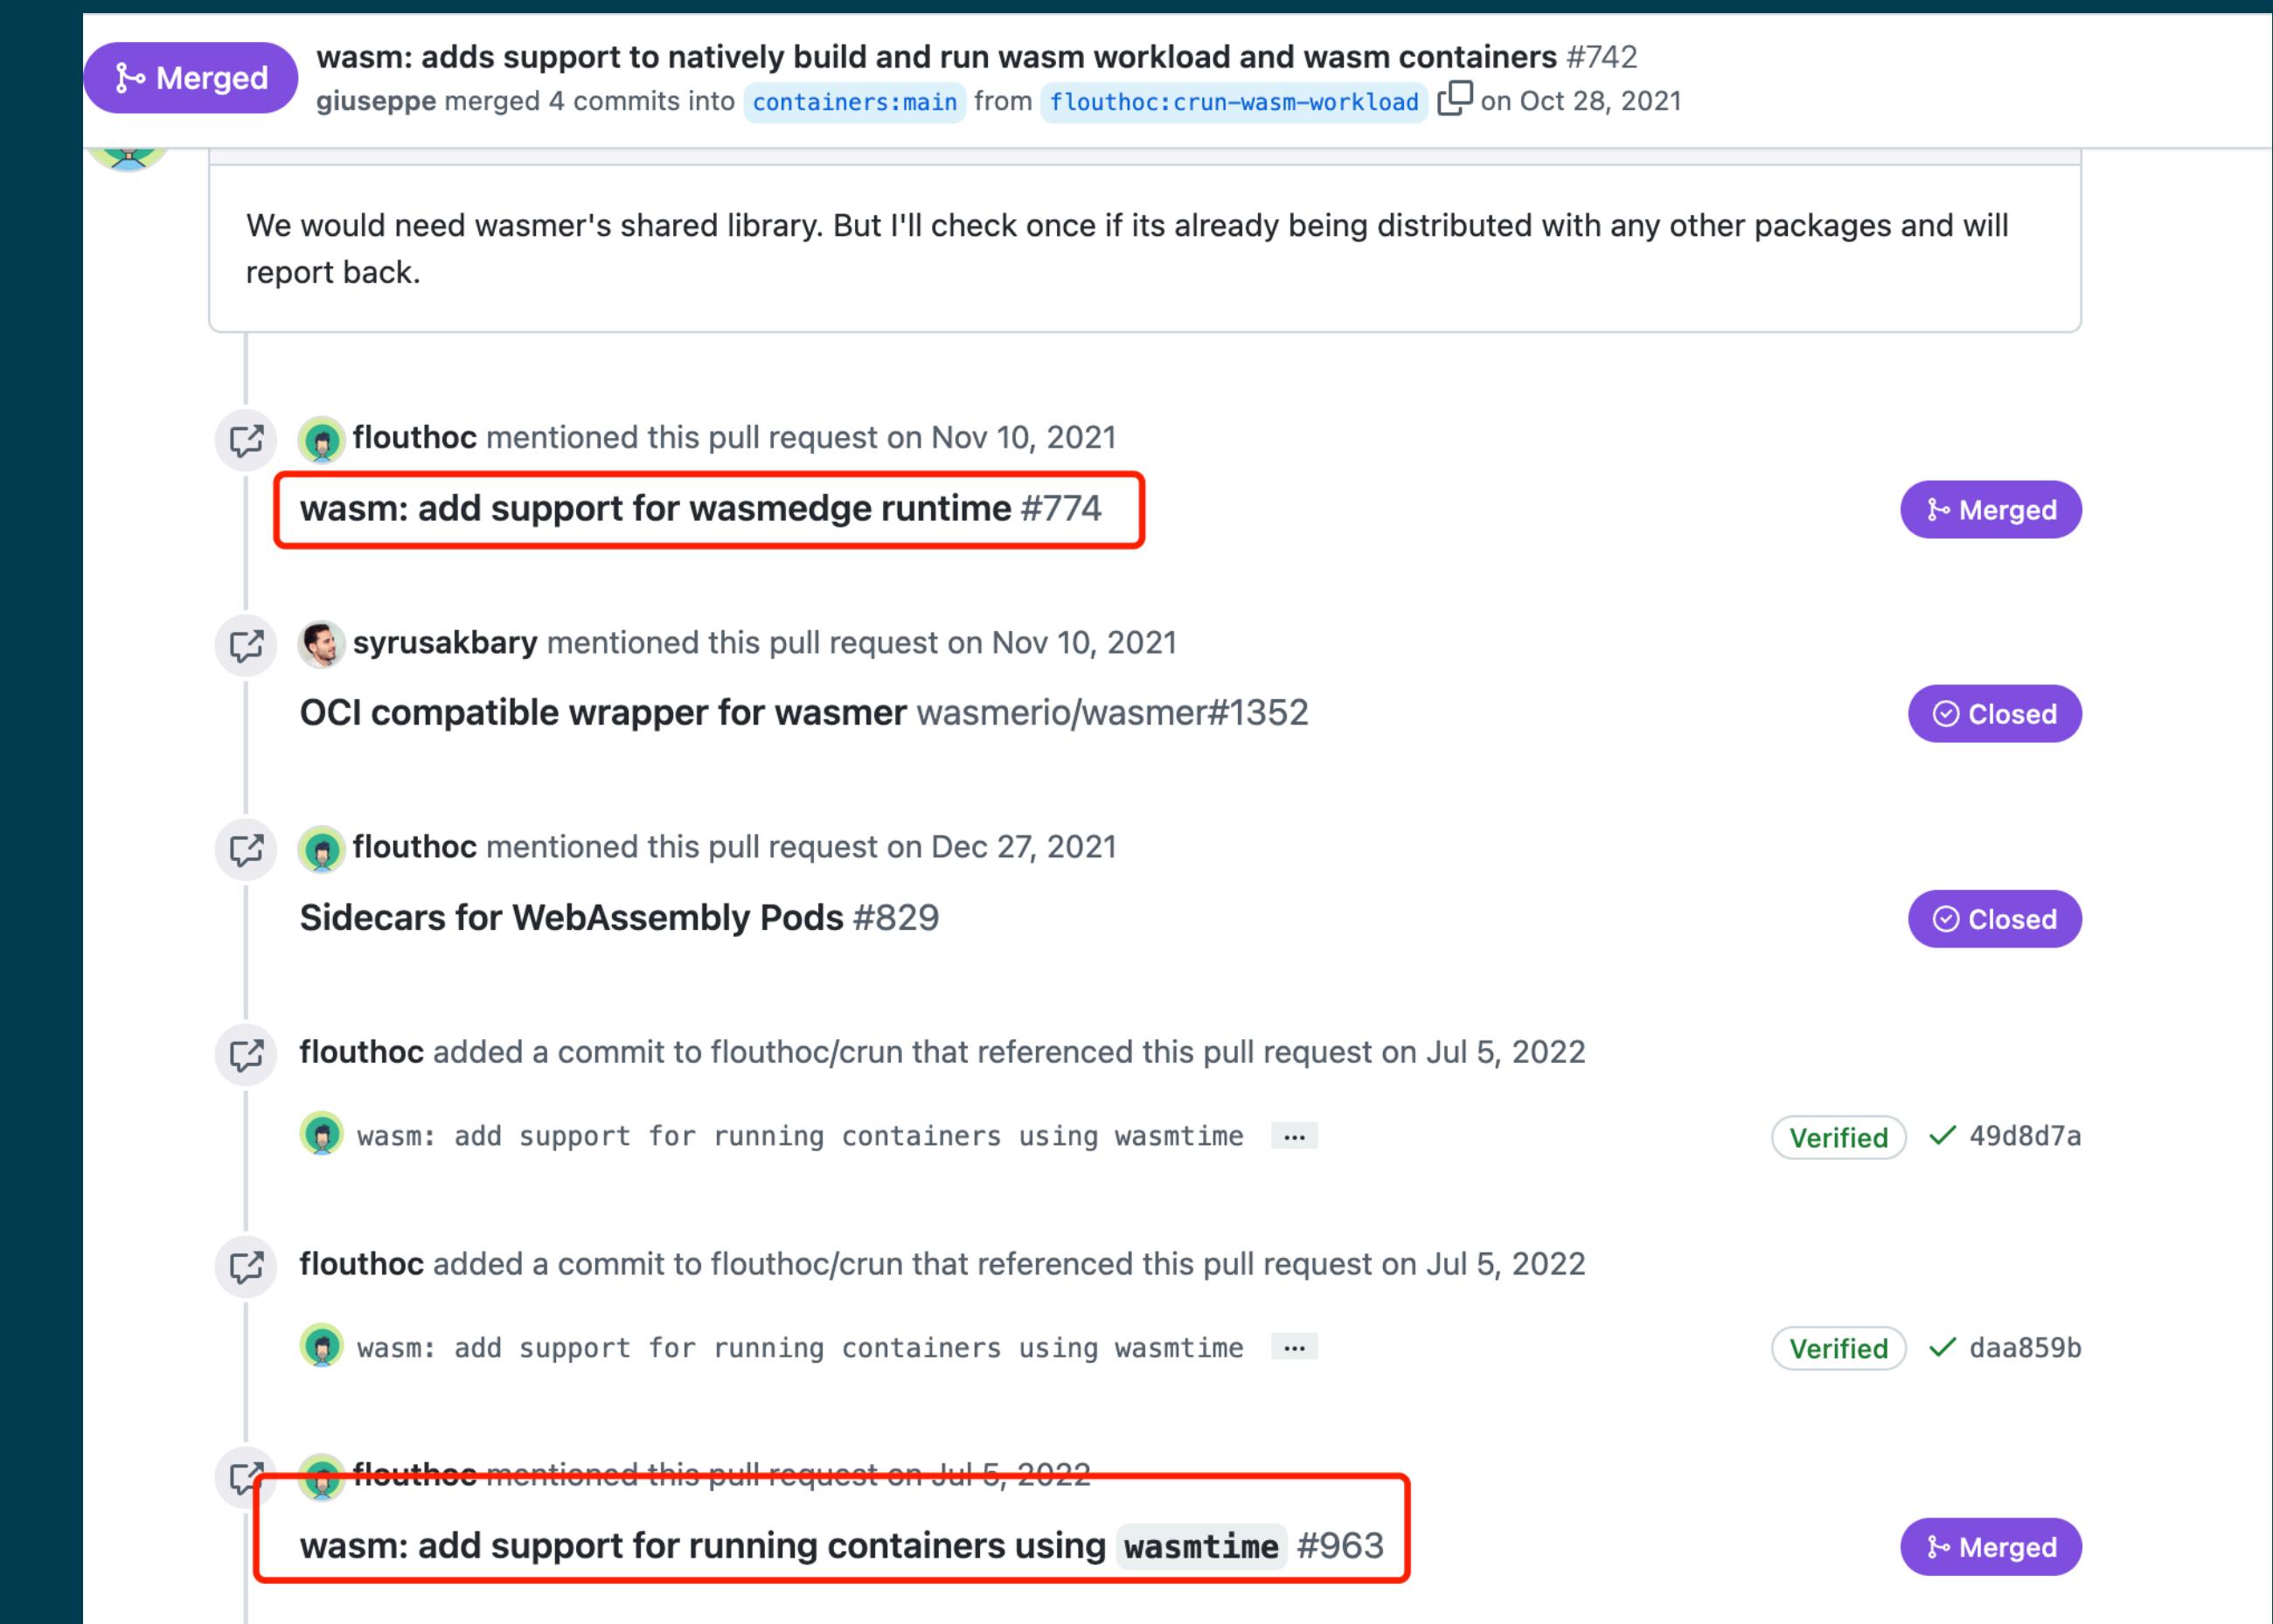
在CRun中运行WASM Workload

CRun为标准的OCI实现

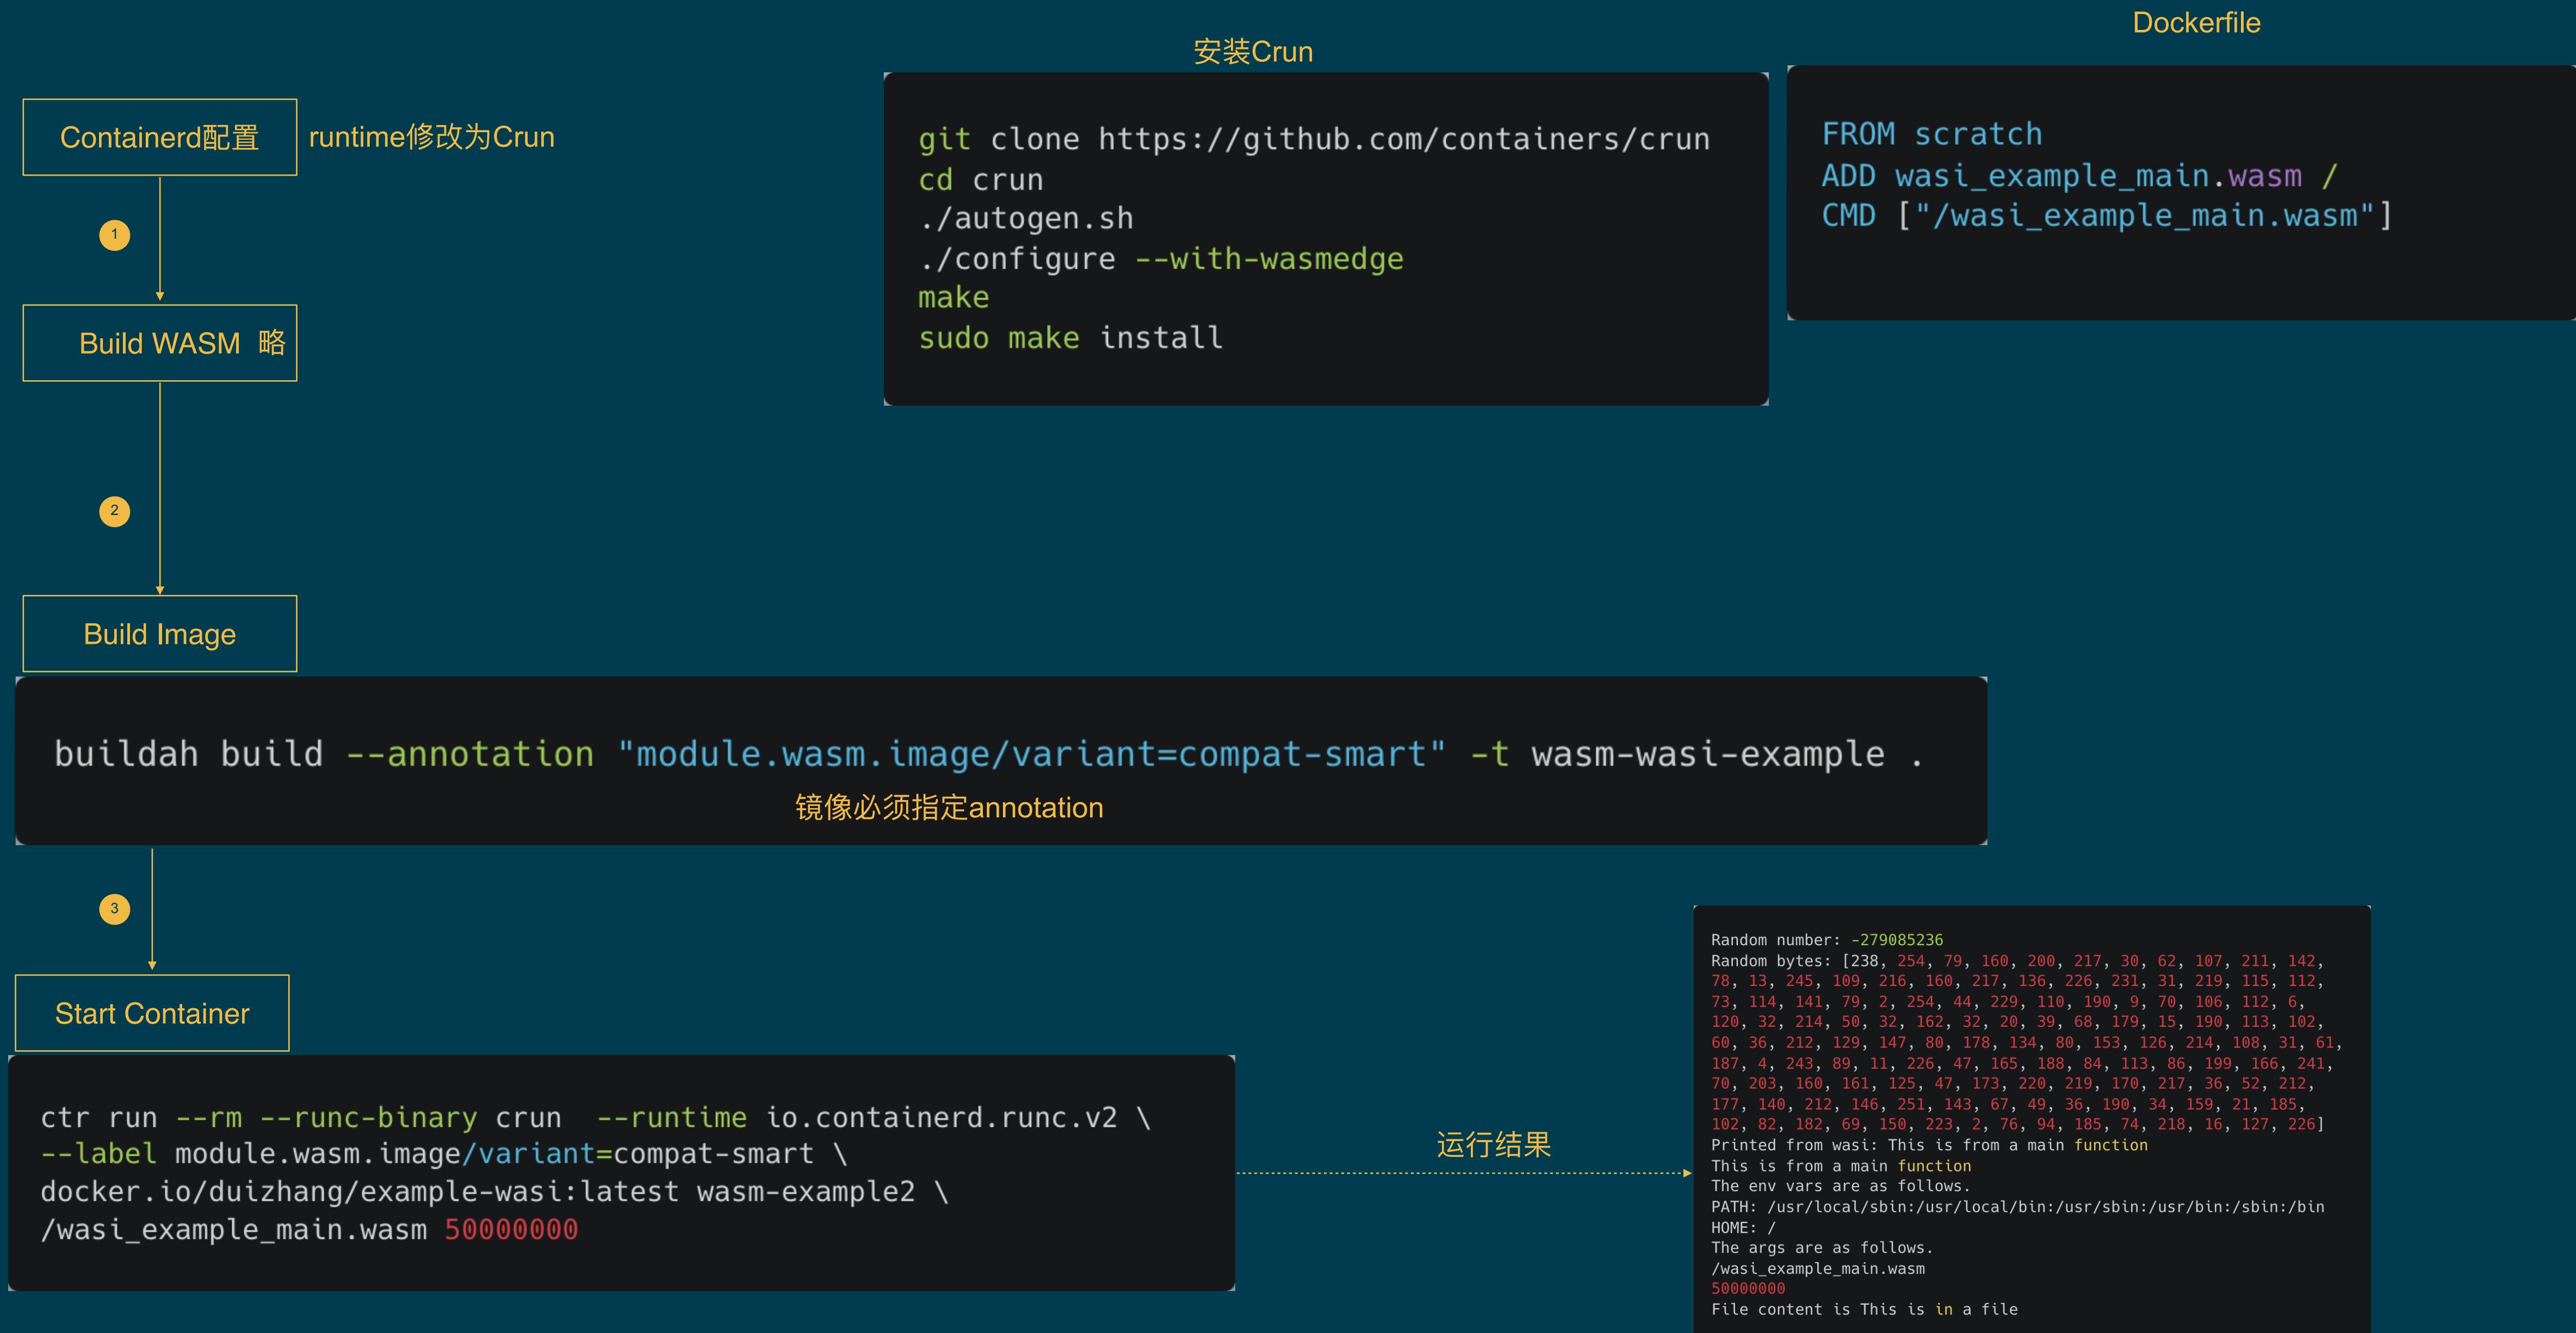
Wasmer

Wasmtime

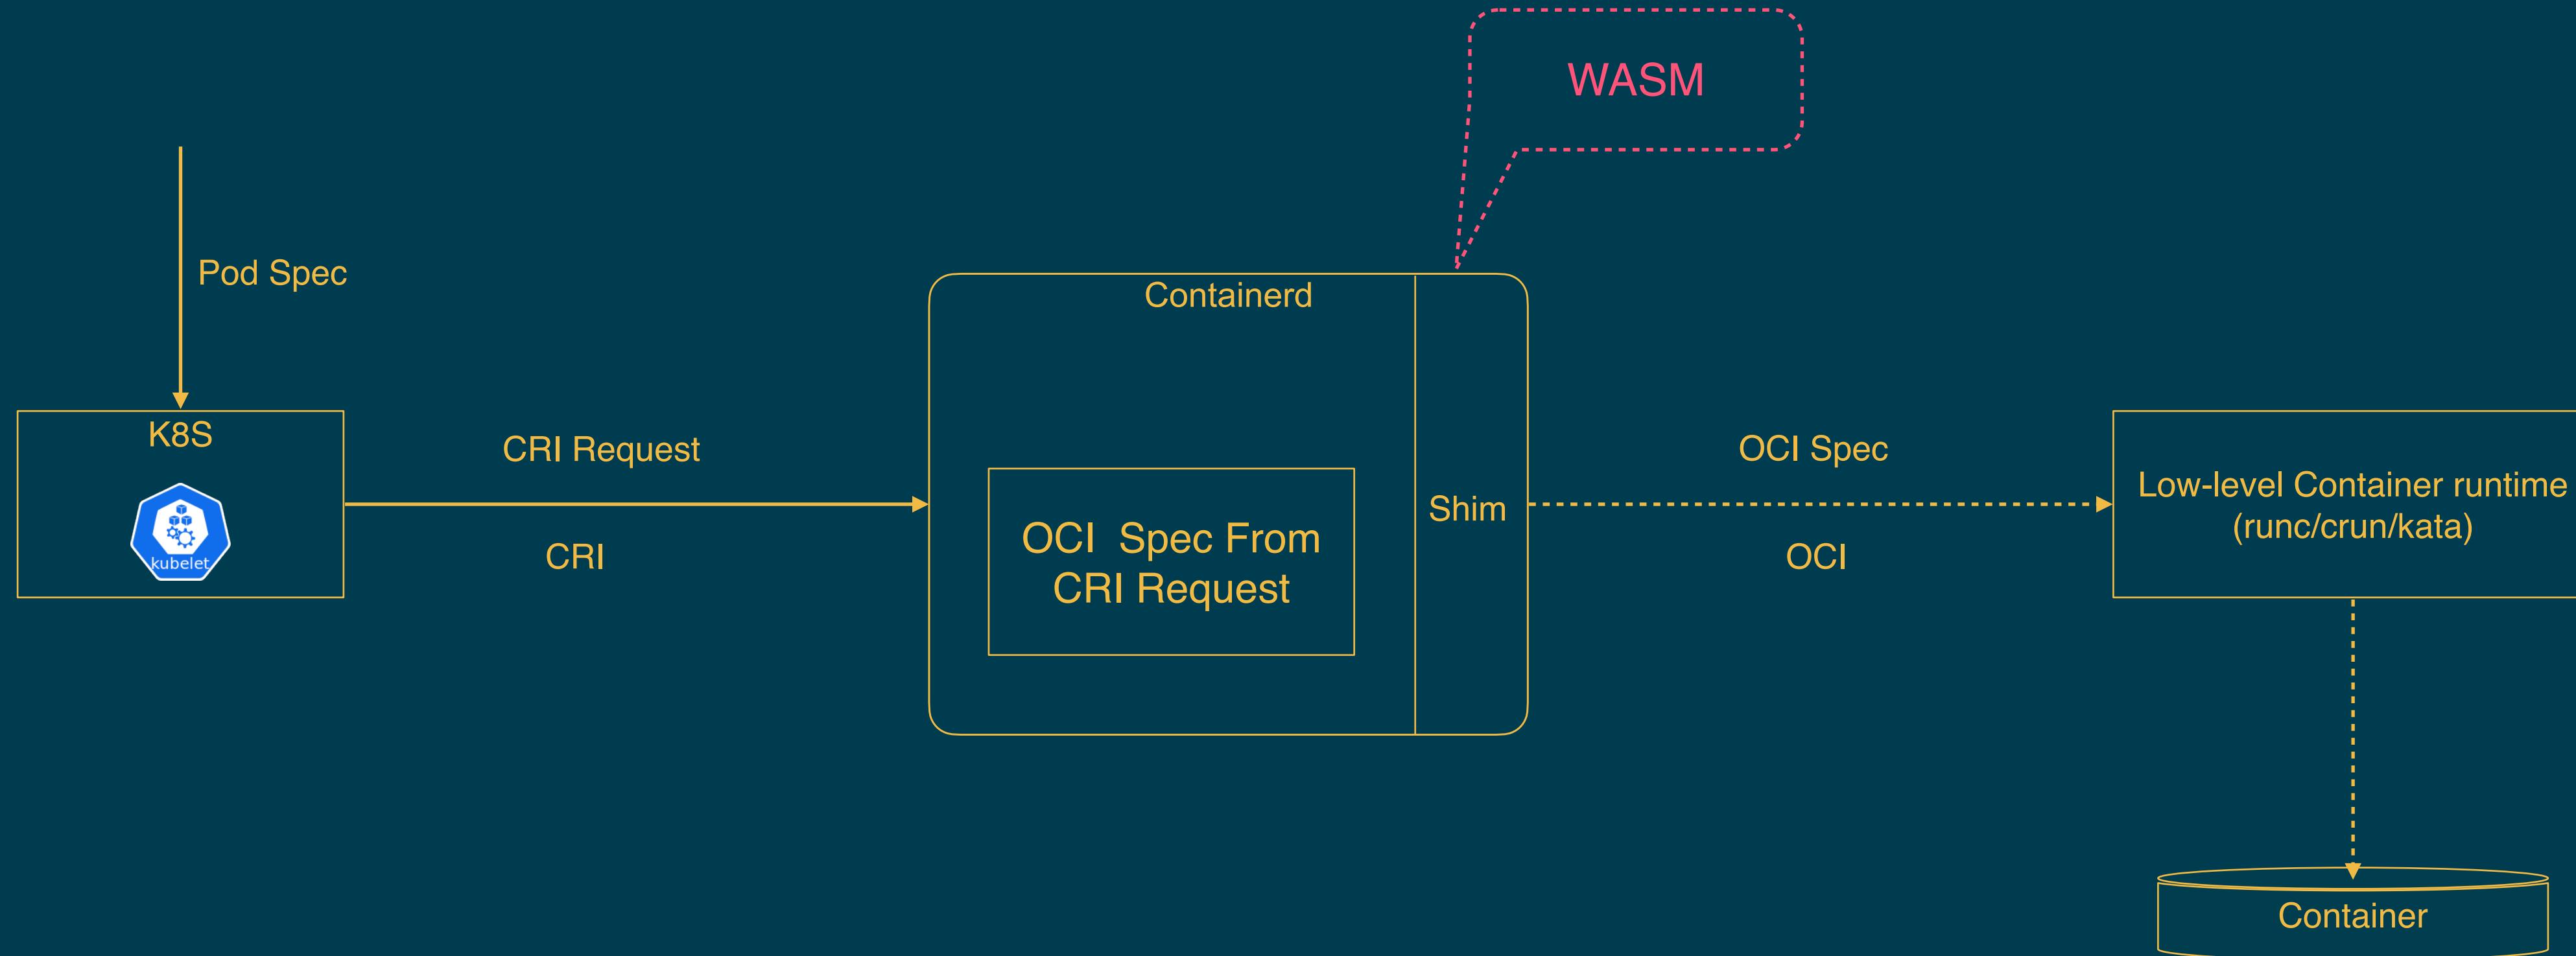
WasmEdge



在CRun中运行WASM Workload

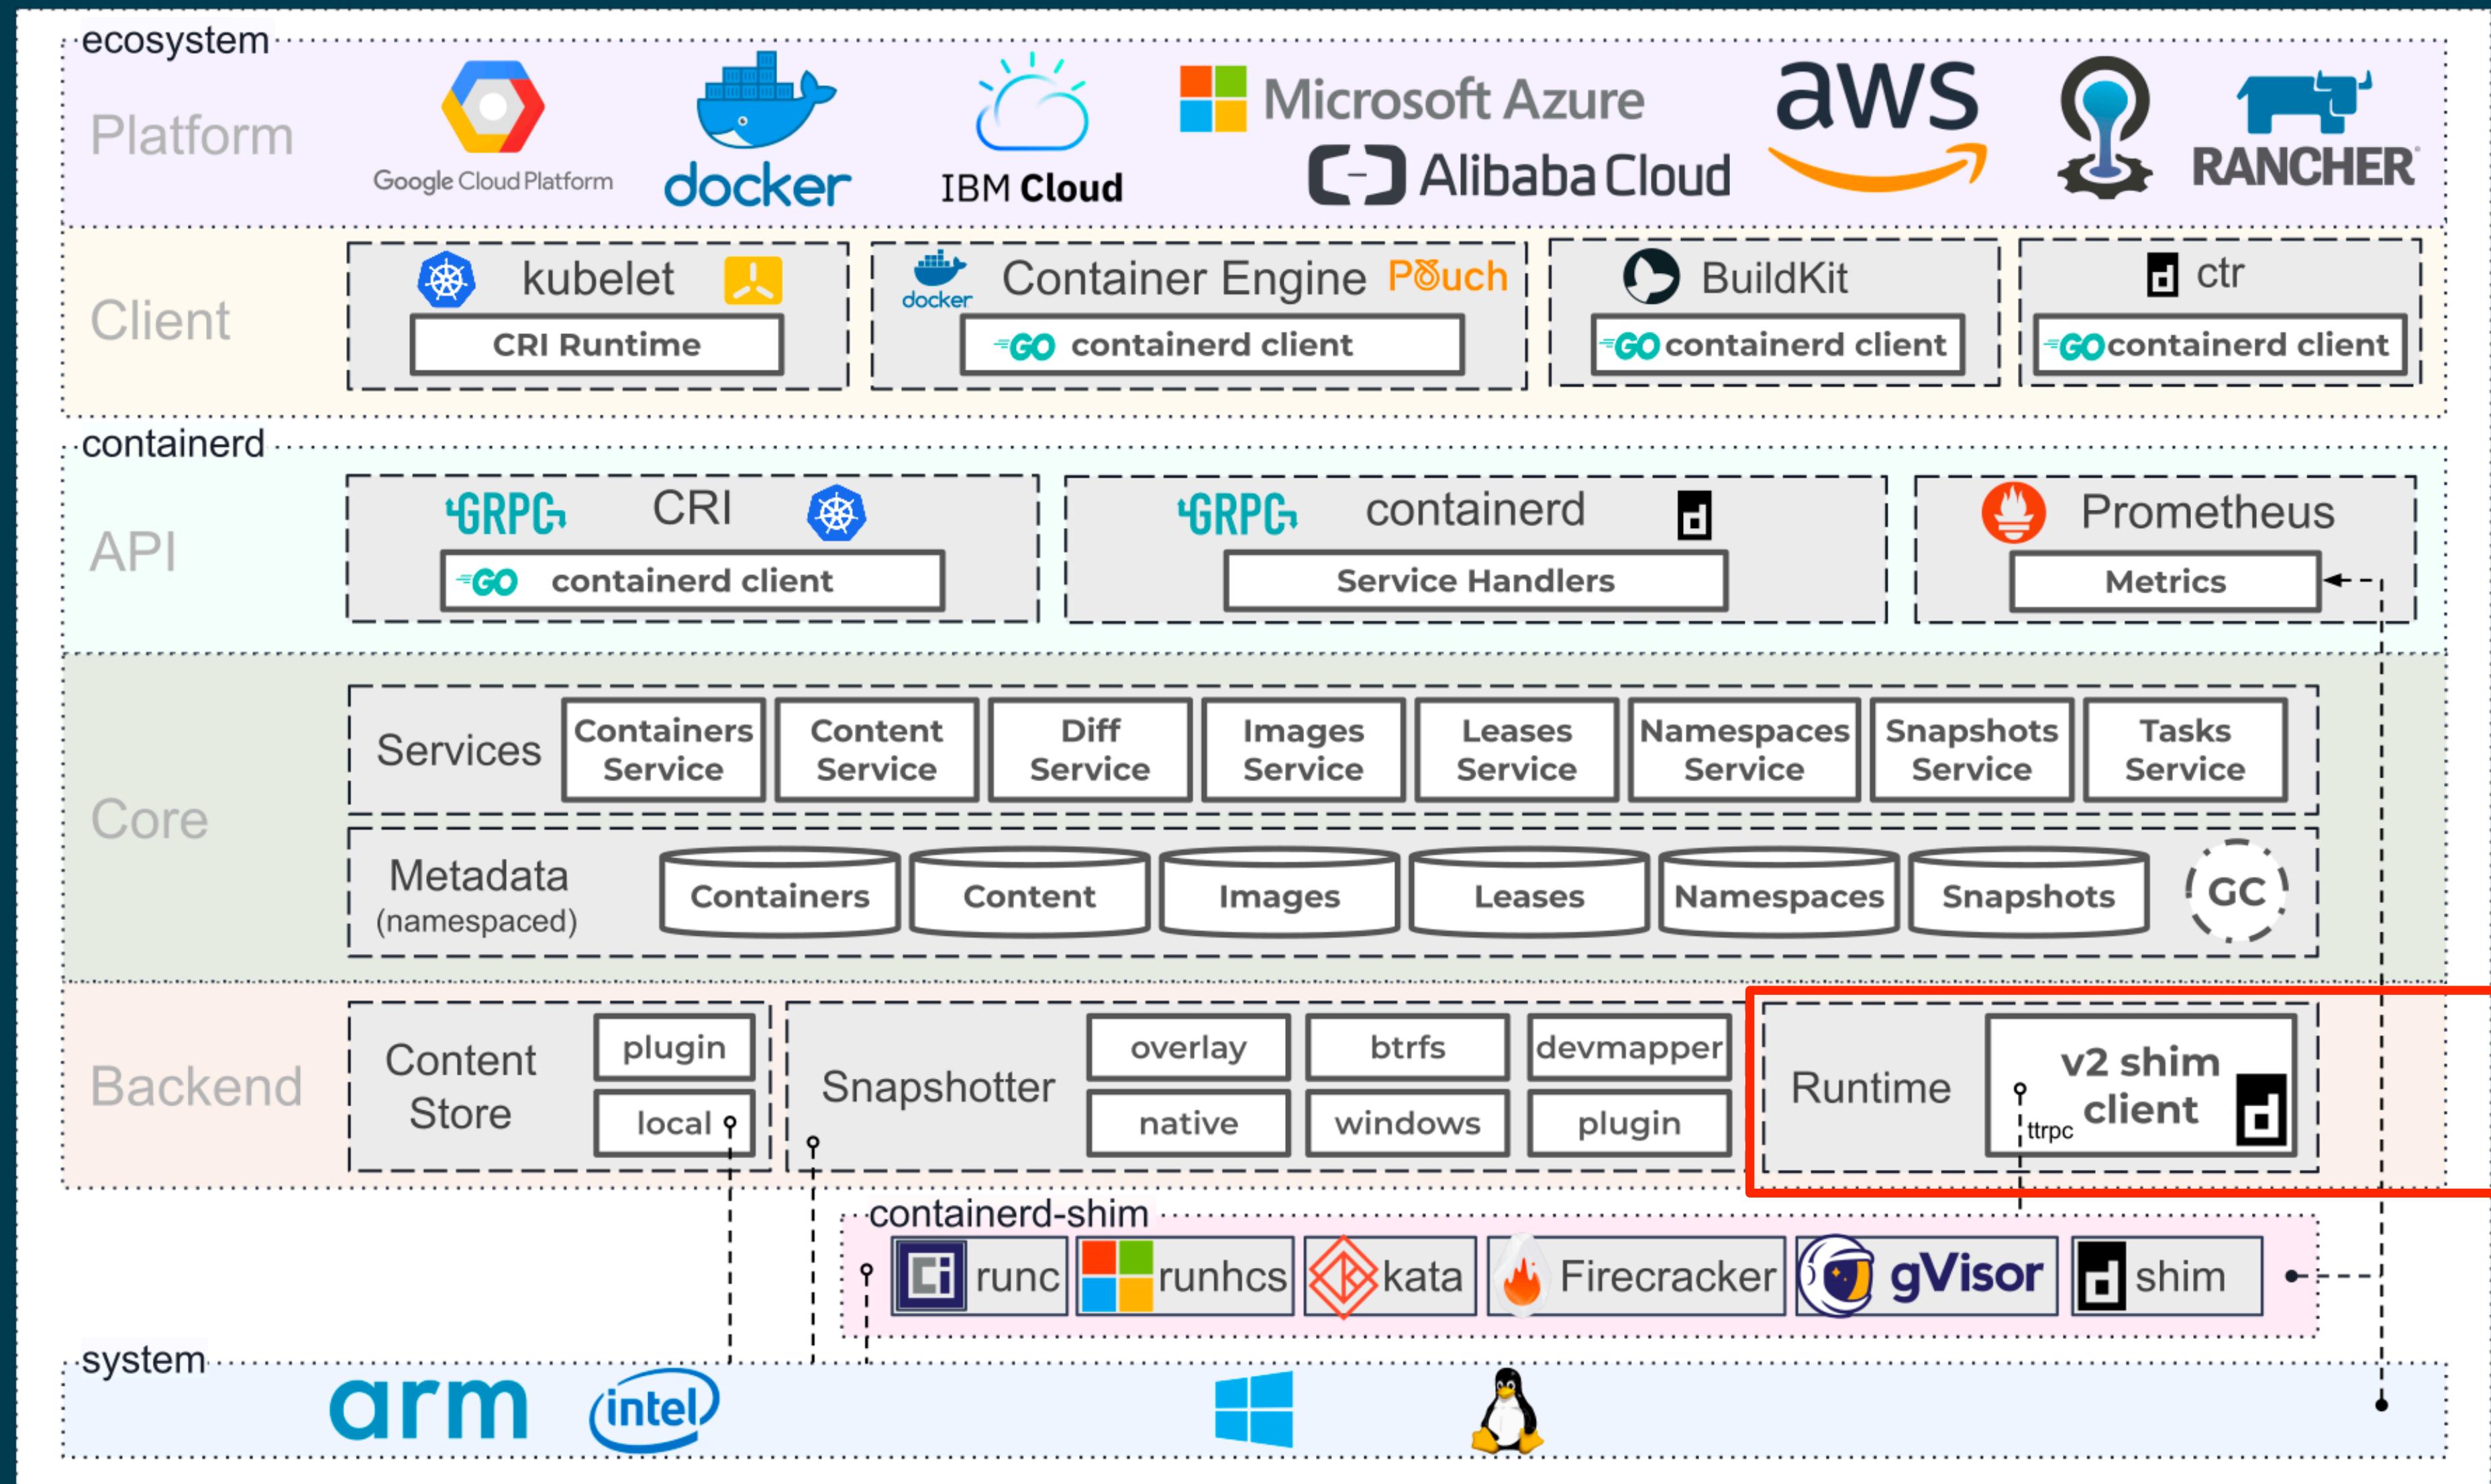


在High-level Container Runtime中支持WASM (Containerd-Shim)



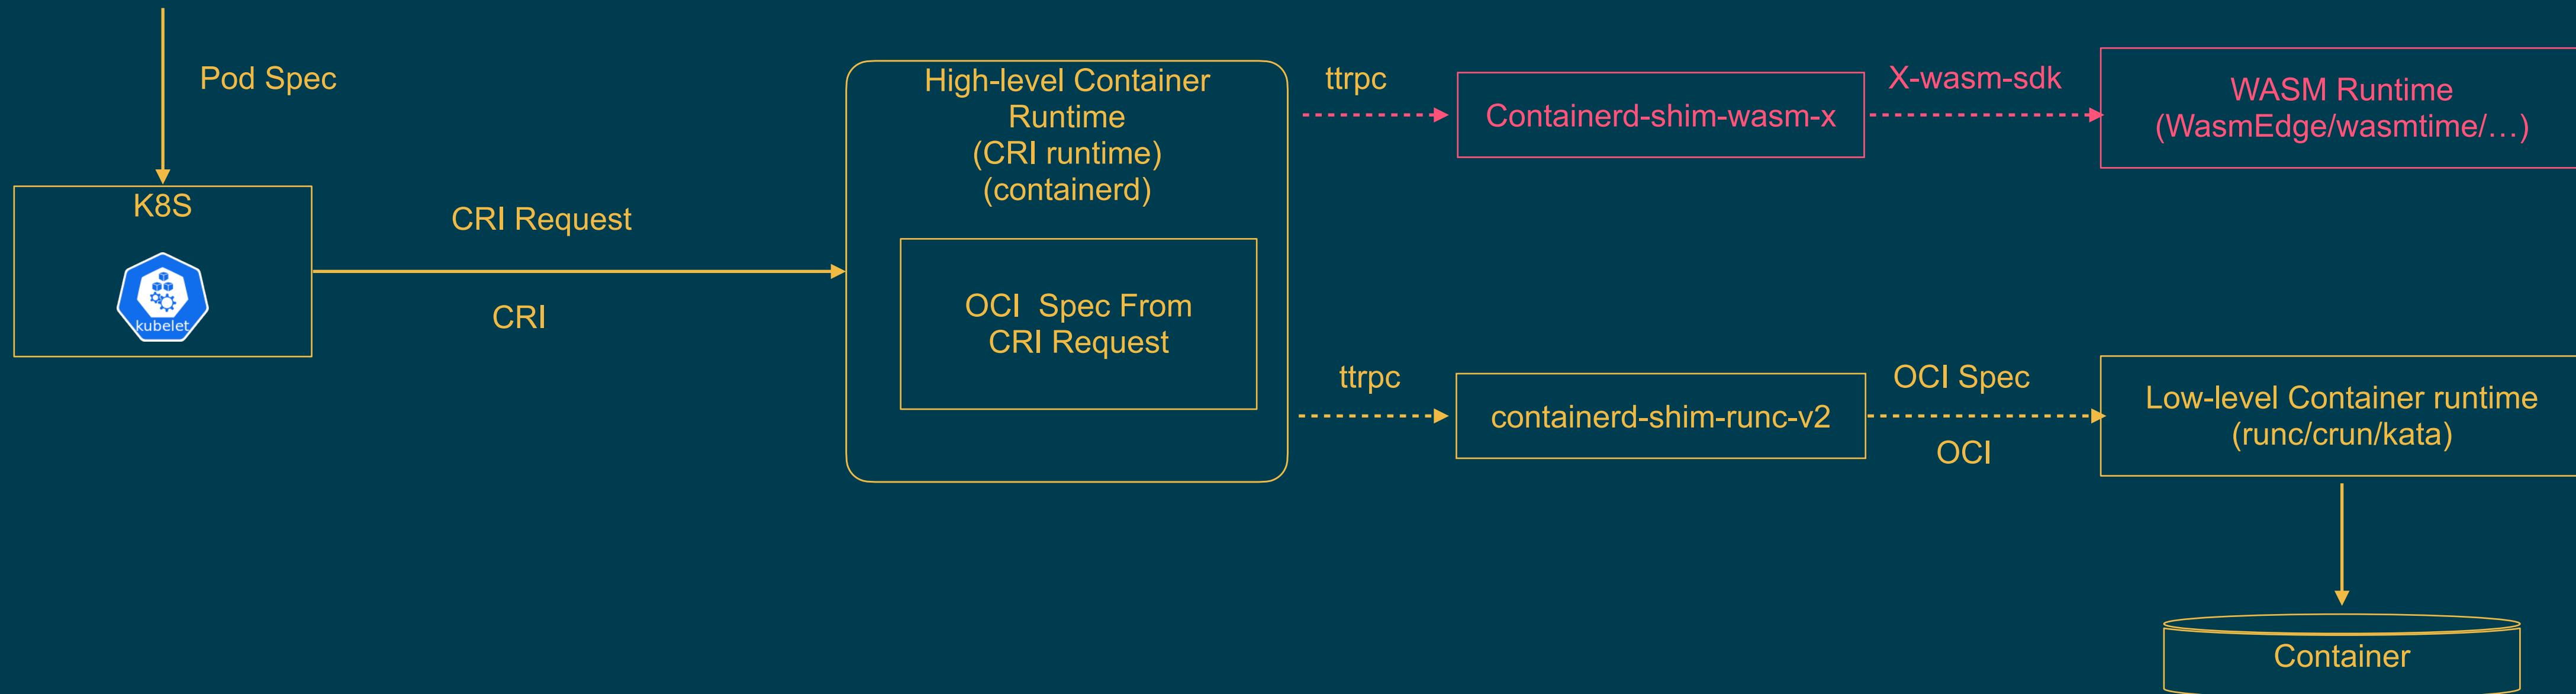
通过新增contianerd-shim支持WASM workload

Contained



Containerd整体架构

Containerd创建WASM Workload流程



通过新增containerd-shim支持WASM Workload

Containerd Shim 接口定义

Shim Server Interface(Golang)

```
type Shim interface {
    shimapi.TaskService
    Cleanup(ctx context.Context) (*shimapi.DeleteResponse, error)
    StartShim(ctx context.Context, opts Start0pts) (string, error)
}

type TaskService interface {
    State(context.Context, *StateRequest) (*StateResponse, error)
    Create(context.Context, *CreateTaskRequest) (*CreateTaskResponse, error)
    Start(context.Context, *StartRequest) (*StartResponse, error)
    Delete(context.Context, *DeleteRequest) (*DeleteResponse, error)
    Pids(context.Context, *PidsRequest) (*PidsResponse, error)
    Pause(context.Context, *PauseRequest) (*emptypb.Empty, error)
    Resume(context.Context, *ResumeRequest) (*emptypb.Empty, error)
    Checkpoint(context.Context, *CheckpointTaskRequest) (*emptypb.Empty, error)
    Kill(context.Context, *KillRequest) (*emptypb.Empty, error)
    Exec(context.Context, *ExecProcessRequest) (*emptypb.Empty, error)
    ResizePty(context.Context, *ResizePtyRequest) (*emptypb.Empty, error)
    CloseIO(context.Context, *CloseIORequest) (*emptypb.Empty, error)
    Update(context.Context, *UpdateTaskRequest) (*emptypb.Empty, error)
    Wait(context.Context, *WaitRequest) (*WaitResponse, error)
    Stats(context.Context, *StatsRequest) (*StatsResponse, error)
    Connect(context.Context, *ConnectRequest) (*ConnectResponse, error)
    Shutdown(context.Context, *ShutdownRequest) (*emptypb.Empty, error)
}
```

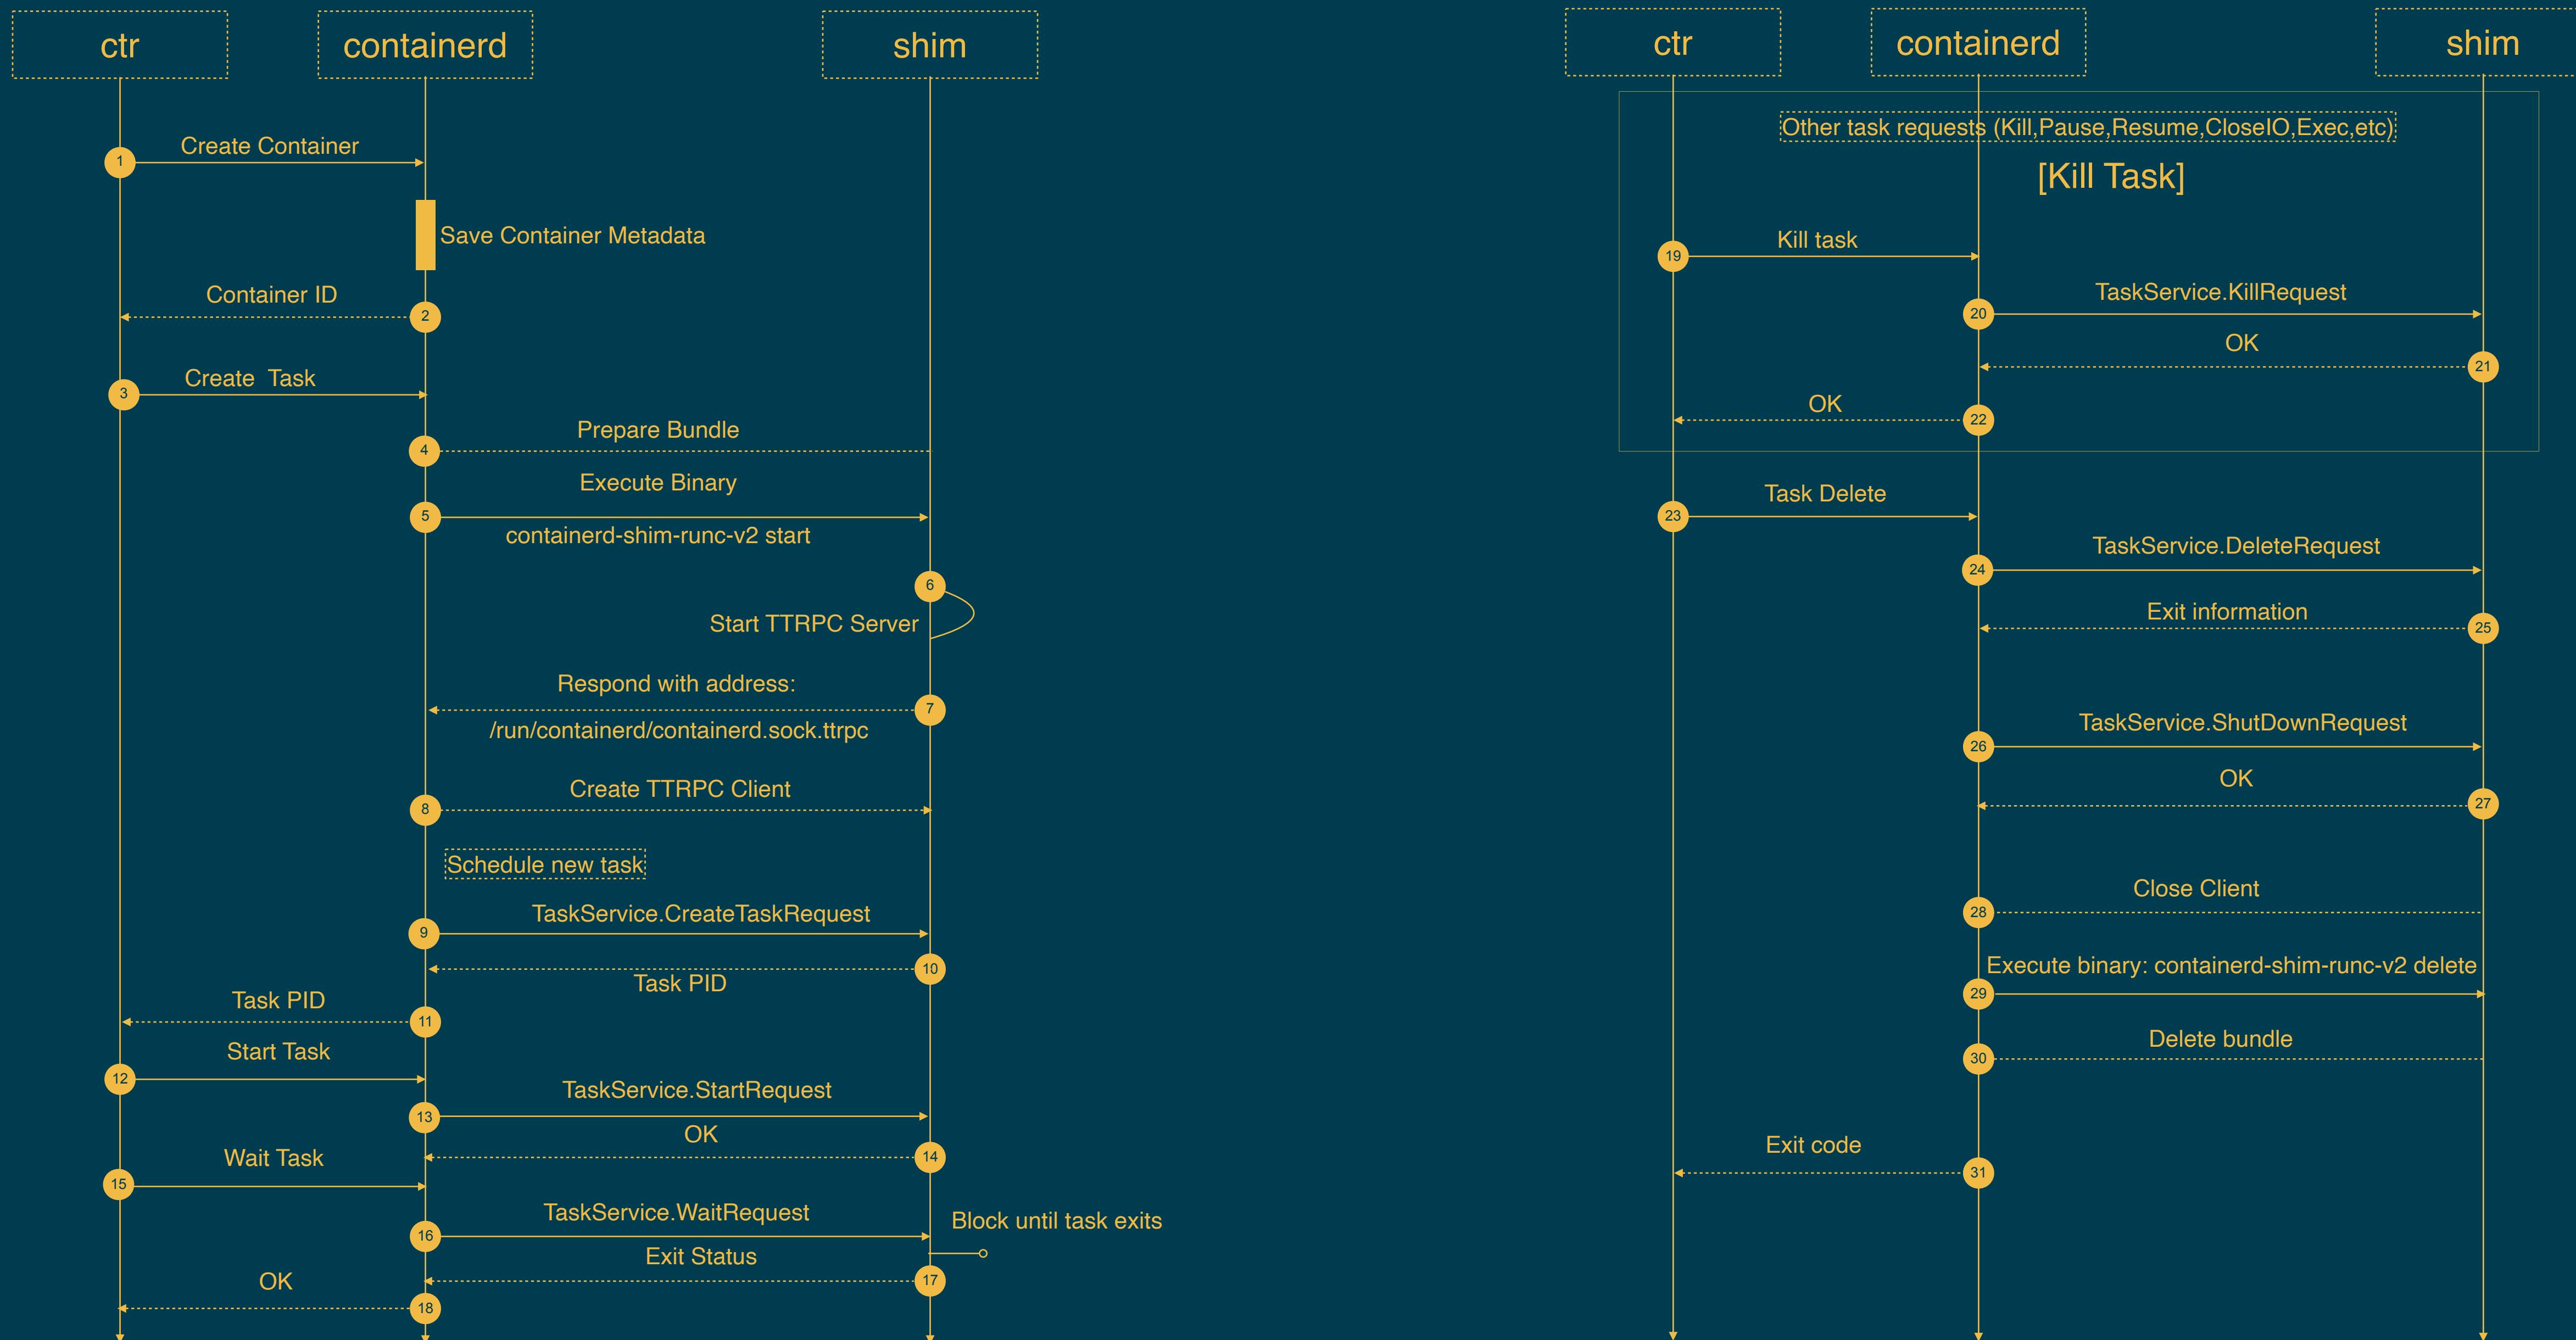
Containerd内置Shim

io.containerd.runc.v1
io.containerd.runc.v2
io.containerd.runhcs.v1(for windows)

三方Shim

containerd-shim-kata-v2(kata-containers)
containerd-shim-runsc-v1(gVisor)
.....

Containererd Shim Flow



Containerd runwasi

runwasi(rust)

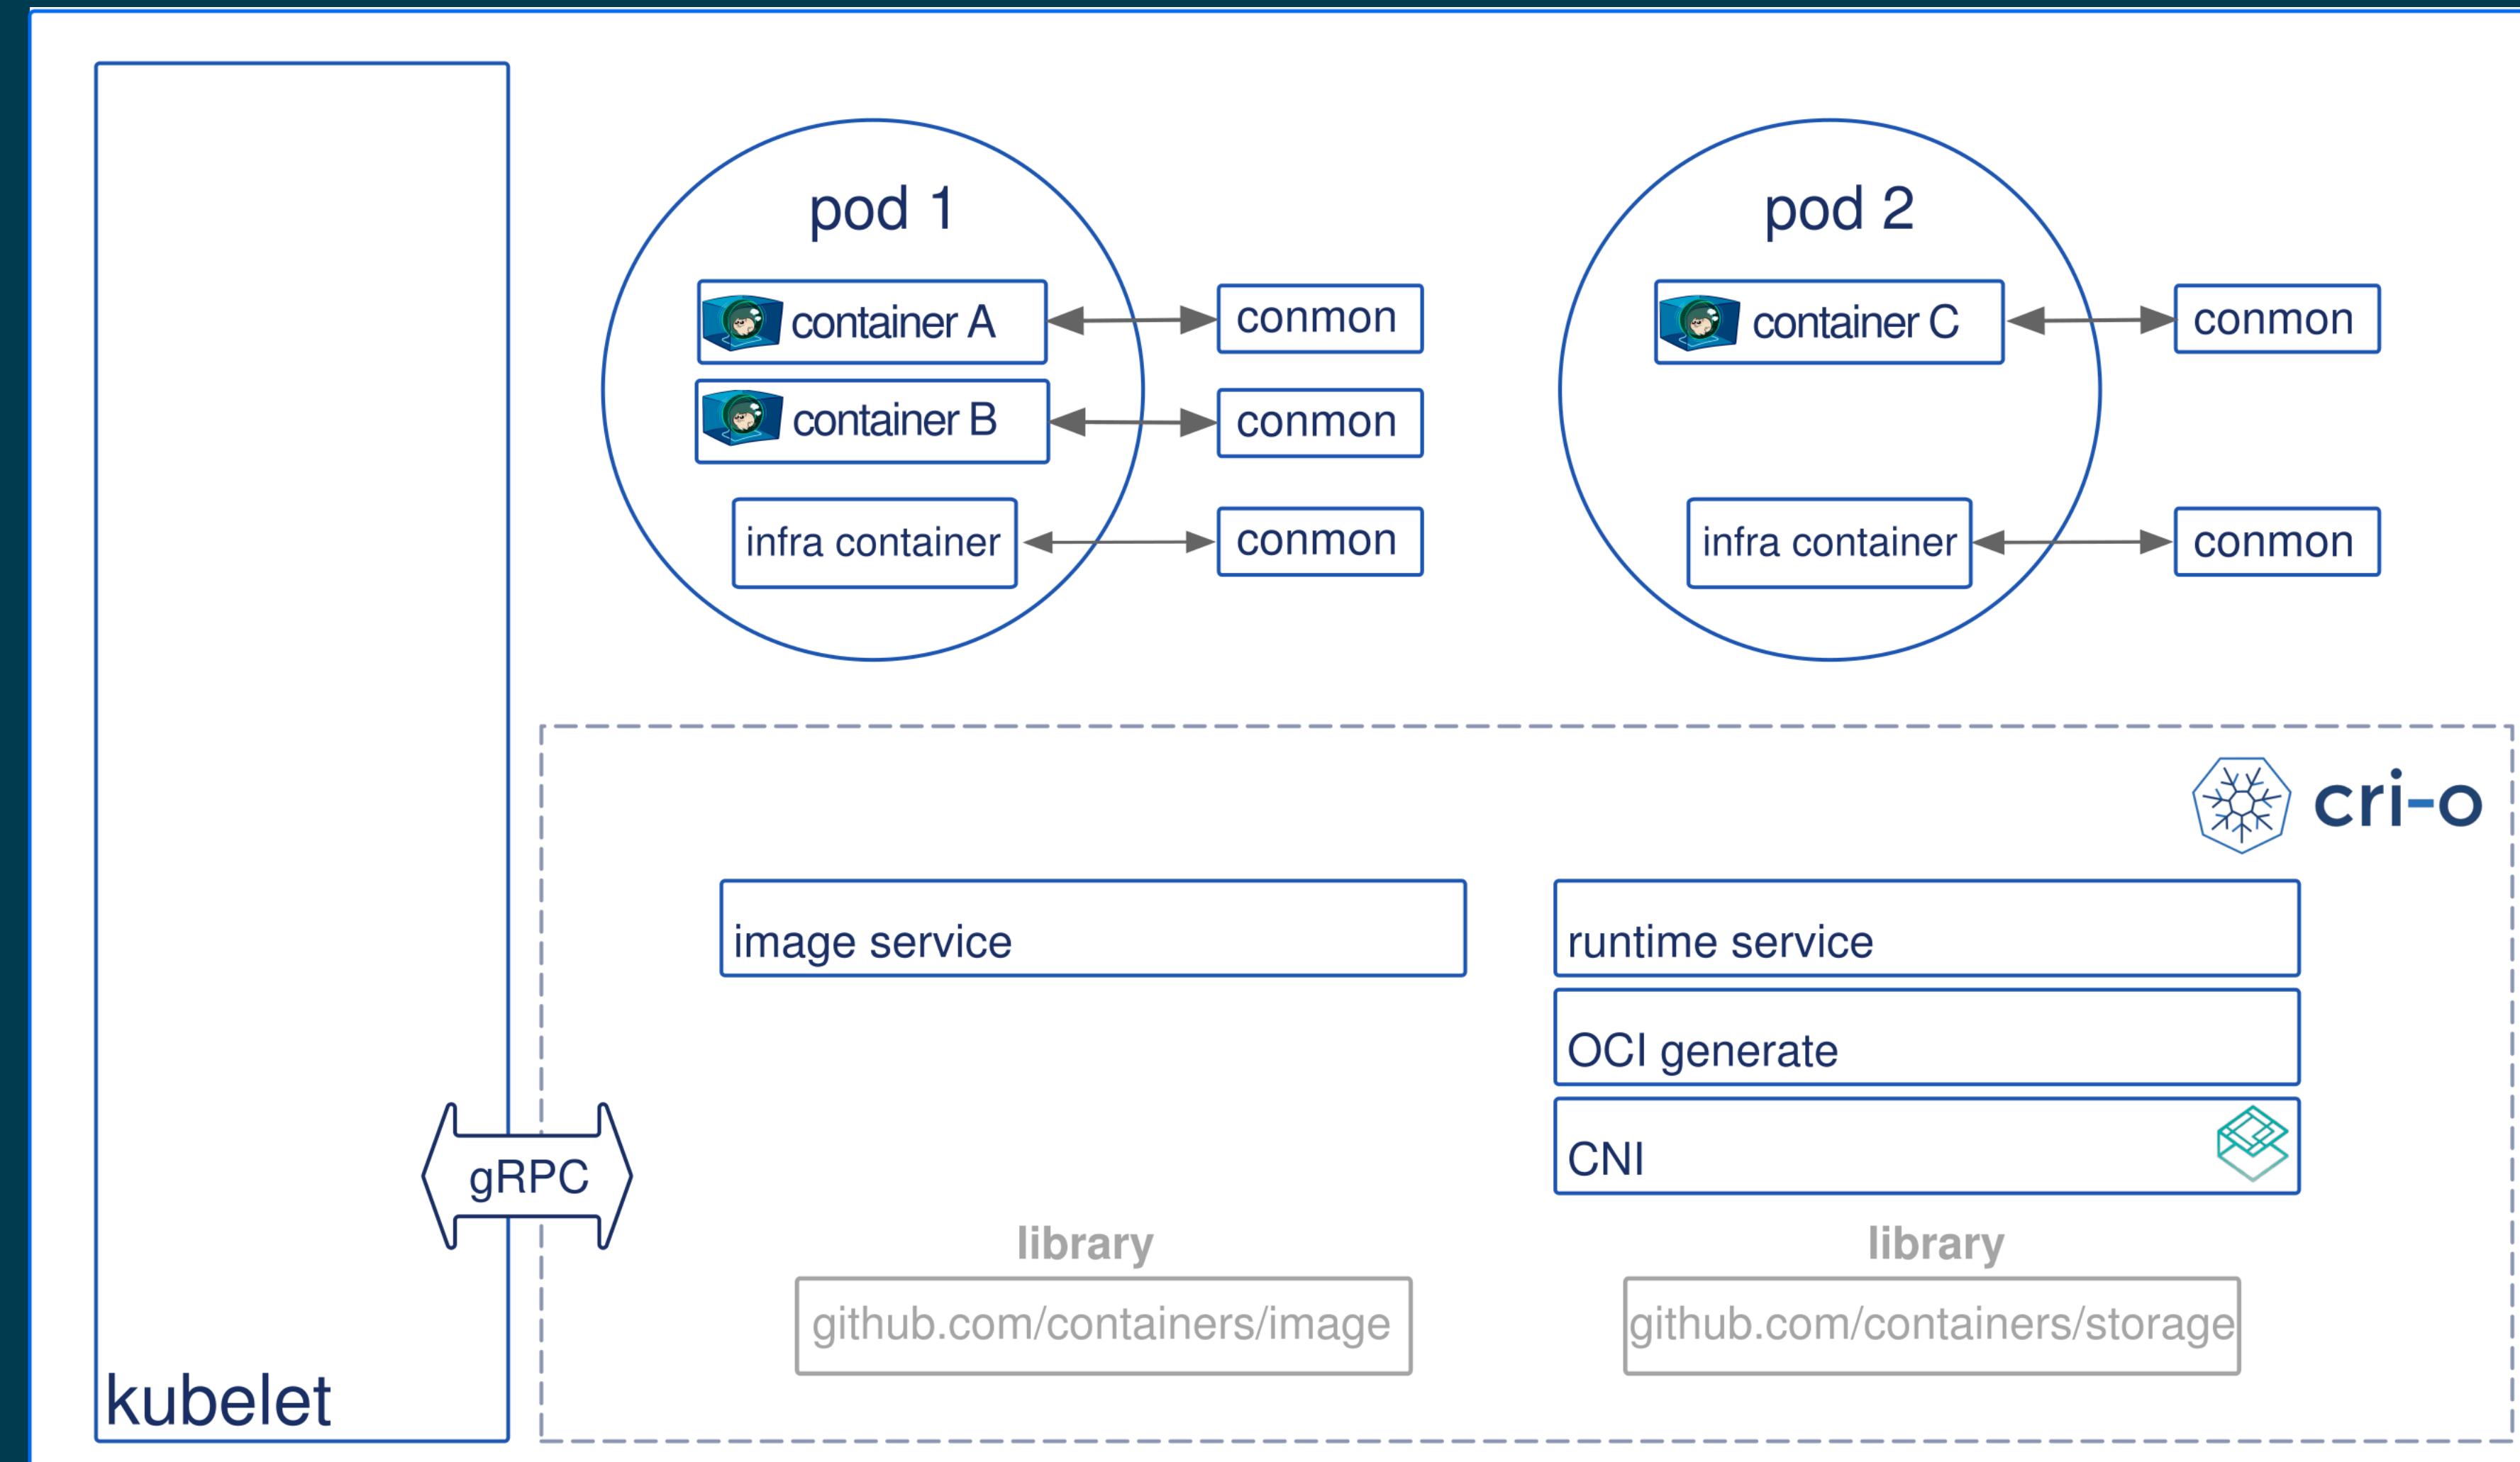
```
pub trait Instance {
    // Create a new instance
    fn new(id: String, cfg: Option<&InstanceConfig<Self::E>>) -> Self;
    // Start the instance and return the pid
    fn start(&self) -> Result<u32, Error>;
    // Send the specified signal to the instance
    fn kill(&self, signal: u32) -> Result<(), Error>;
    // Delete the instance
    fn delete(&self) -> Result<(), Error>;
    // wait for the instance to exit and send the exit code and exit timestamp to the provided sender.
    fn wait(&self, send: Sender<(u32, DateTime<Utc>)>) -> Result<(), Error>;
}
```

Containerd Shim实例—wasmtime

以containerd-shim-wasmtime-v1为例

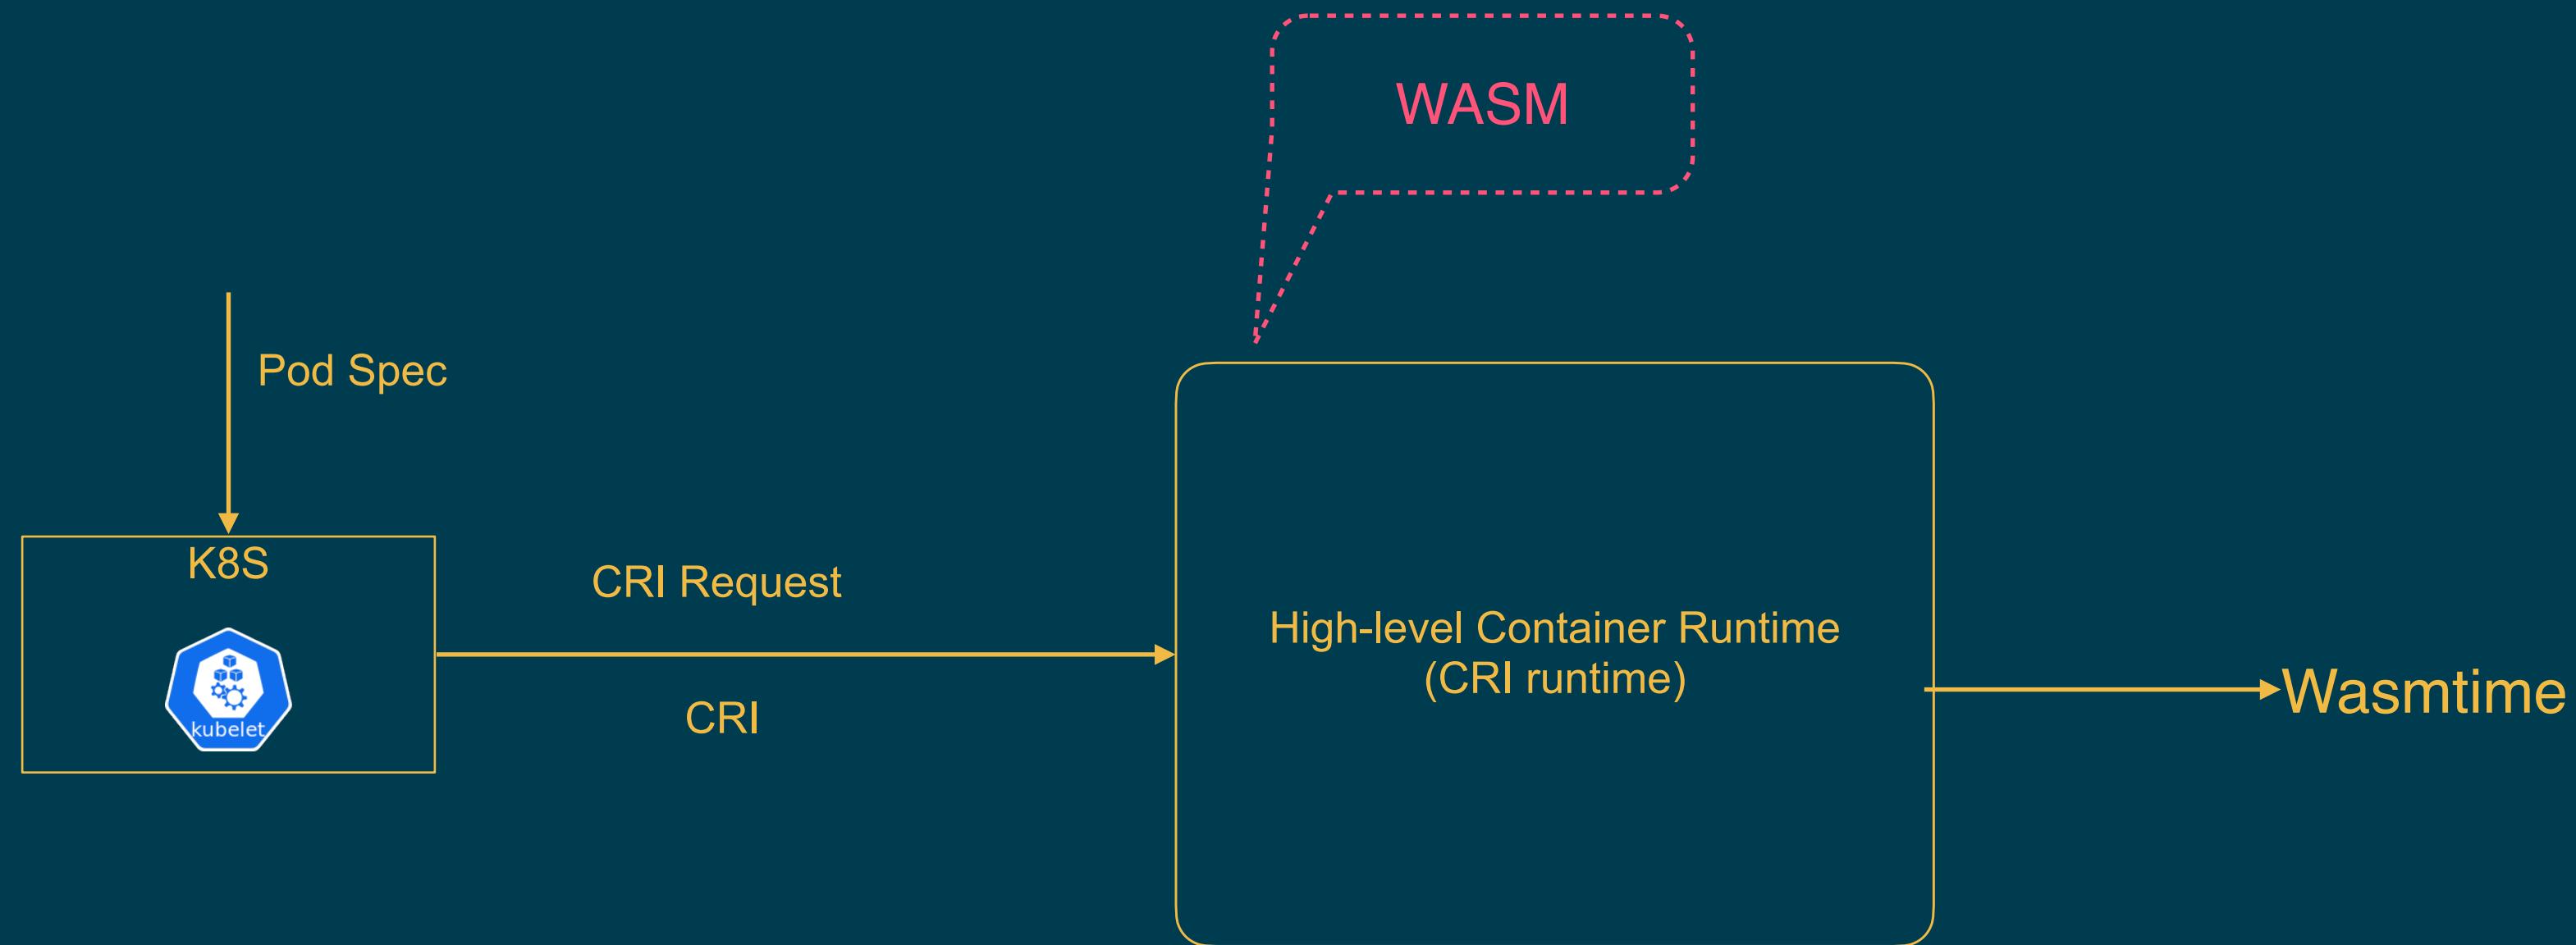
```
ctr run --rm --runtime=io.containerd.wasmtime.v1 \
ghcr.io/containerd/runwasi/wasi-demo-app:latest \
testwasm2 /wasi-demo-app.wasm echo 'hello'
```

CRI-O



每一个container都由一个单独的common进程管理

实现新的CRI组件支持WASM Workload



通过实现新的CRI支持WASM Workload

WASMRun为例

① kubelet配置中指定runtime为wasmrun

```
KUBELET_KUBEADM_ARGS="--container-runtime-endpoint=unix:///var/run/cri-wasm.sock --pod-infra-container-image=registry.k8s.io/pause:3.9"
```

② 查看节点的CONTAINER-RUNTIME

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE	KERNEL-VERSION	CONTAINER-RUNTIME
vm-0-10-ubuntu	NotReady	<none>	70d	v1.26.3	172.19.0.10	<none>	Ubuntu 20.04.6 LTS	5.4.0-144-generic	wasmrun://0.1.0
vm-0-11-ubuntu	Ready	control-plane	74d	v1.26.0	172.19.0.11	<none>	Ubuntu 20.04 LTS	5.4.0-126-generic	containerd://1.6.8

③ 创建测试用Job，并查看Pod日志

```
zwliu@liuzhenweideMacBook-Pro ~ % kubectl get node -o wide
NAME           STATUS      ROLES   AGE    VERSION
vm-0-10-ubuntu NotReady   <none>   70d   v1.26.3
vm-0-11-ubuntu Ready      control-plane   74d   v1.26.0
zwliu@liuzhenweideMacBook-Pro KubeFunction % kubectl get pod
NAME          READY   STATUS    RESTARTS   AGE
wasmrun-job-mg4h6   0/1     Completed   0          6s
wasmtime-5f5645c98c-2gqwz   1/1     Running    0          15h
zwliu@liuzhenweideMacBook-Pro KubeFunction % kubectl logs -f wasmrun-job-mg4h6
Hello, duizhangwww!
```

```
go run main.go -wasm-runtime-endpoint unix:///var/run/cri-wasm.sock \
    -container-runtime-socket=/run/containerd/containerd.sock \
    -v 3
```

wasmrun运行参数

```
apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: "wasmrun"
handler: "wasmrun"
scheduling:
  nodeSelector:
    "app": "wasmrun"
```

runtimeClass定义

```
apiVersion: batch/v1
kind: Job
metadata:
  name: wasmrun-job
spec:
  template:
    metadata:
      labels:
        app: wasmrun
    spec:
      runtimeClassName: wasmrun
      restartPolicy: Never
      containers:
        - name: testwasm
          image: docker.io/duizhang/hello-wasm:9.0
          command: ["greet"]
          args: ["duizhangwww"]
```

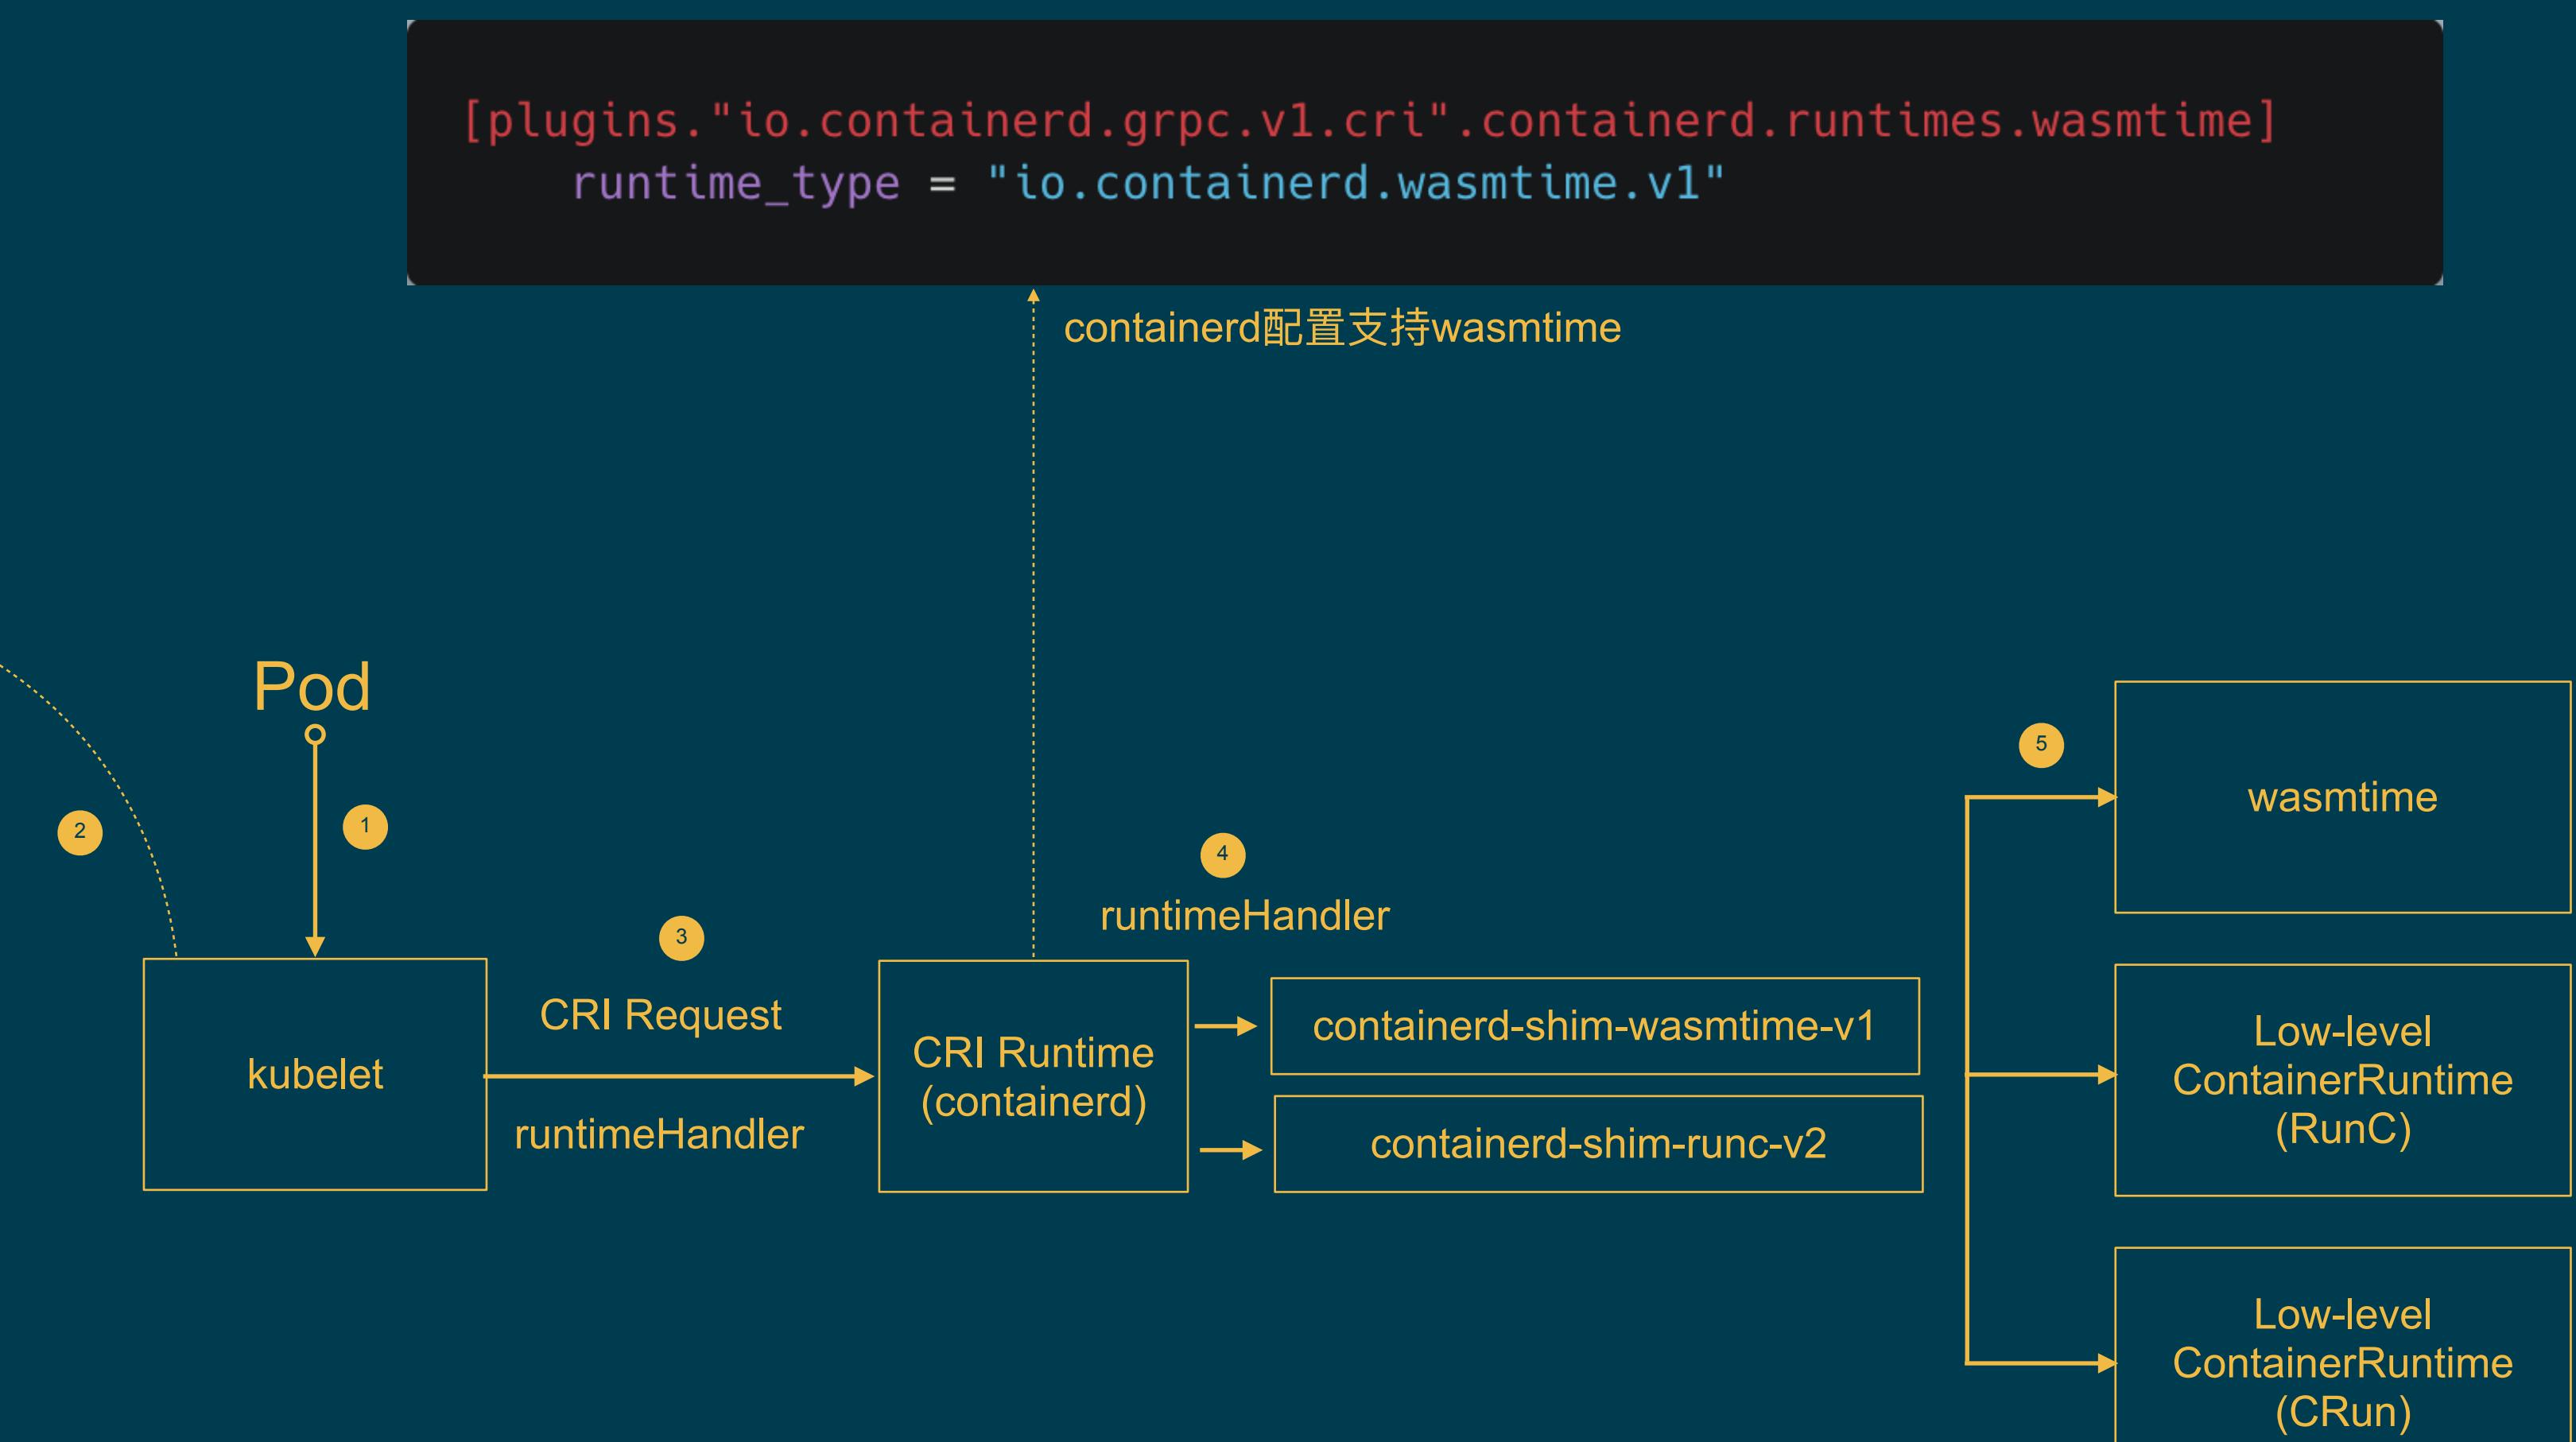
测试用Job

Kubernetes多运行时整体流程

```
apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: "wasmtime"
handler: "wasmtime"
scheduling:
  nodeSelector:
    "wasm": "wasmtime"
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wasmtime
spec:
  replicas: 1
  selector:
    matchLabels:
      app: wasmtime
  template:
    metadata:
      labels:
        app: wasmtime
  spec:
    runtimeClassName: wasmtime
    containers:
    - name: wasm-hello
      image: docker.io/duizhang/wasi-demo-app:latest
      command: ["/wasi-demo-app.wasm", "echo", "duizhang"]
```

测试用工作负载



```
zwliu@liuzhenweideMacBook-Pro KubeFunction % kubectl get pod
NAME                  READY   STATUS    RESTARTS   AGE
wasmtime-5f5645c98c-2gqwz  1/1     Running   0          9m45s
zwliu@liuzhenweideMacBook-Pro KubeFunction % kubectl logs -f wasmtime-5f5645c98c-2gqwz
duizhang
exiting
结果输出
```

定义新的kubelet以支持WASM Workload

krustlet

```
kubectl get nodes -o wide
NAME     STATUS   ROLES      AGE    VERSION   INTERNAL-IP        EXTERNAL-IP   OS-IMAGE         KERNEL-VERSION   CONTAINER-RUNTIME
minikube Ready    master     18m    v1.18.0   192.168.99.165   <none>       Buildroot 2019.02.10   4.19.107        docker://19.3.8
krustlet Ready    agent      9s     v1.17.0   10.0.2.2        <none>       <unknown>        <unknown>        mvp
```

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-wasm
spec:
  containers:
  - name: hello-wasm
    image: webassembly.azurecr.io/hello-wasm:v1
  tolerations:
  - effect: NoExecute
    key: kubernetes.io/arch
    operator: Equal
    value: wasm32-wasi  # or wasm32-wasmcloud according to module target arch
  - effect: NoSchedule
    key: kubernetes.io/arch
    operator: Equal
    value: wasm32-wasi  # or wasm32-wasmcloud according to module target arch
```

谢谢

姓名 刘振伟

职位

email@thoughtworks.com

/thoughtworks